

# Adversarial Search and Game Playing

(Respect your opponent to make good decisions)

**Russell and Norvig:**  
**Chapter 6 (4th edition)**

Slides adopted from Jean-Claude Latombe at Stanford University  
(used with permission)

# Games



Games like Chess or Go are compact settings that mimic the uncertainty of interacting with the natural world.

For centuries humans have used them to exert their intelligence.

Recently, there has been great success in building game programs that challenge human supremacy.

# Specific Setting

- two-player
- turn-taking:  
at each step only one player chooses a move
- deterministic
- fully observable
- zero-sum:  
the rewards of all players add up to zero  
= if one player wins, the other player loses
- time-constrained game:  
only a limited amount of time to make a move,  
typically not enough to “solve” the game



# Search Problem Formulation

Initial state: Game setup at start

Player(s): Which player moves in state

Action(s): Legal moves in a state

Result(s,a): Transition model

Terminal-Test(s): True when game over

Evaluate(s, p): Estimate of how good s is for player p

# Simplification

For two-player zero-sum games:

- we name the players **MAX** and **MIN**
- **MAX** wants to maximize his score
- **MIN** wants to minimize **MAX**'s score

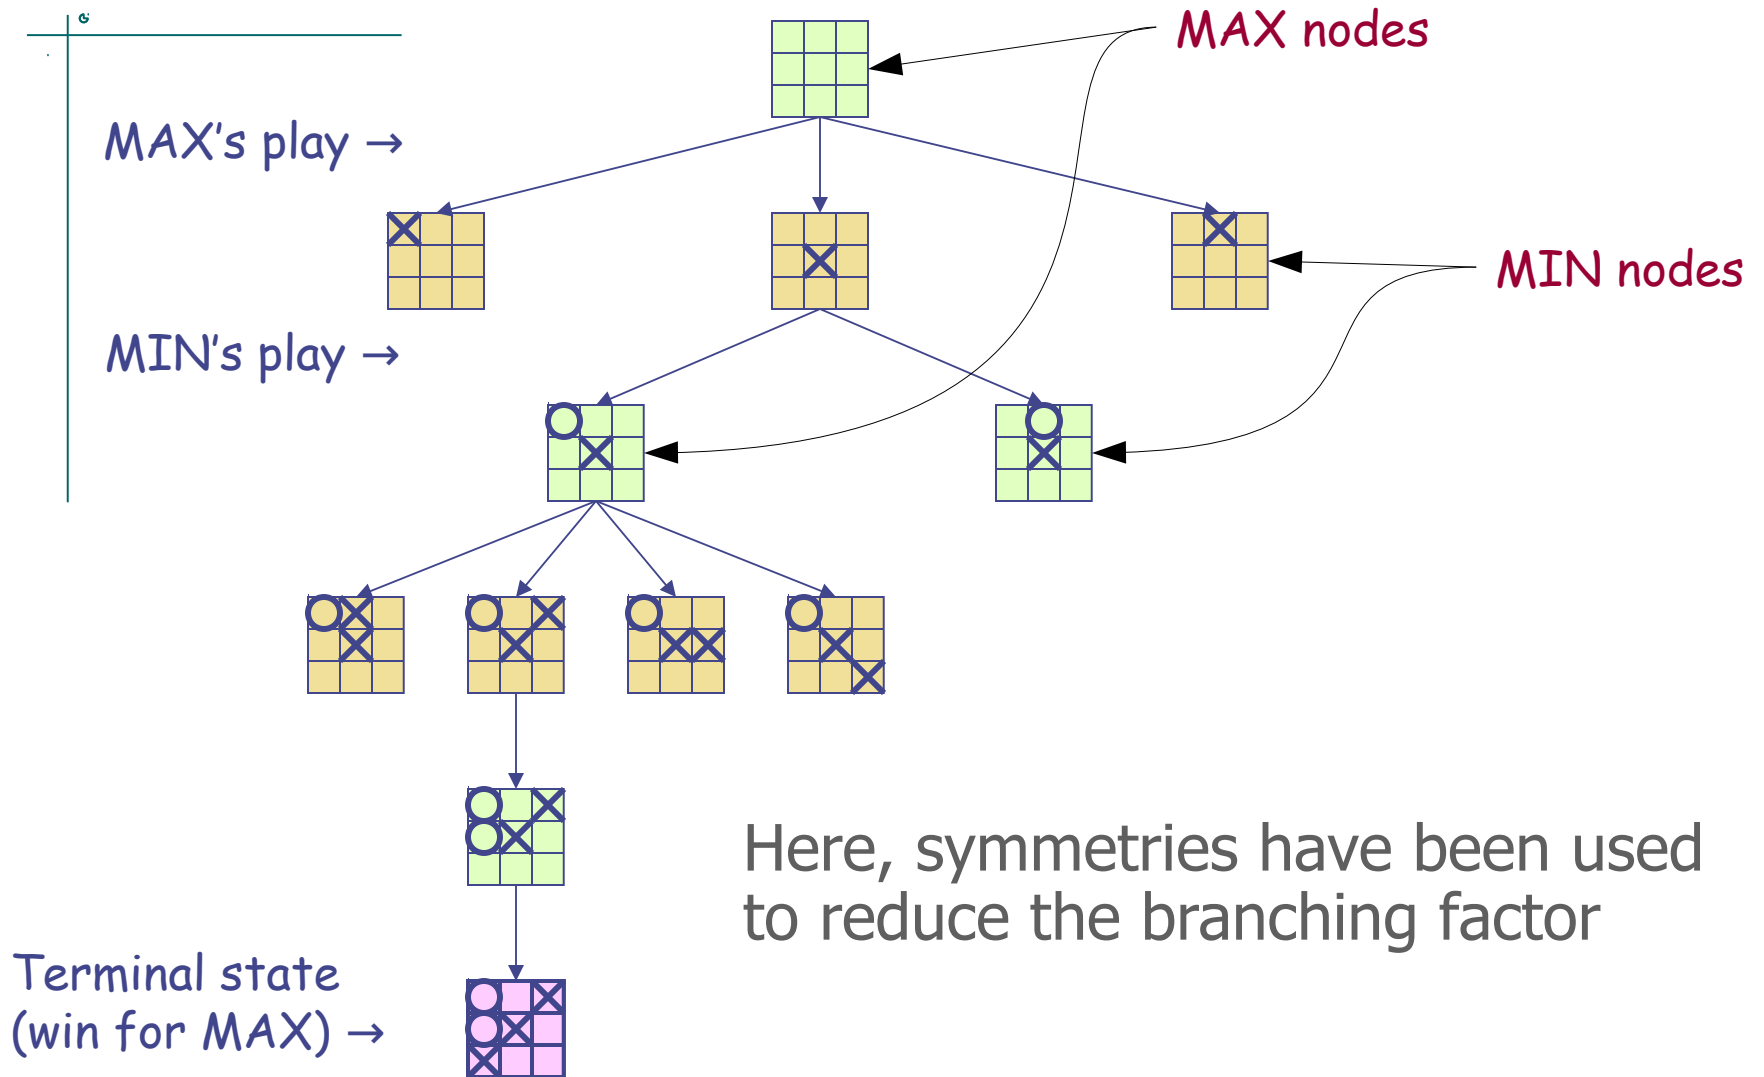
→ we only consider the score of **MAX**

# Time Limit

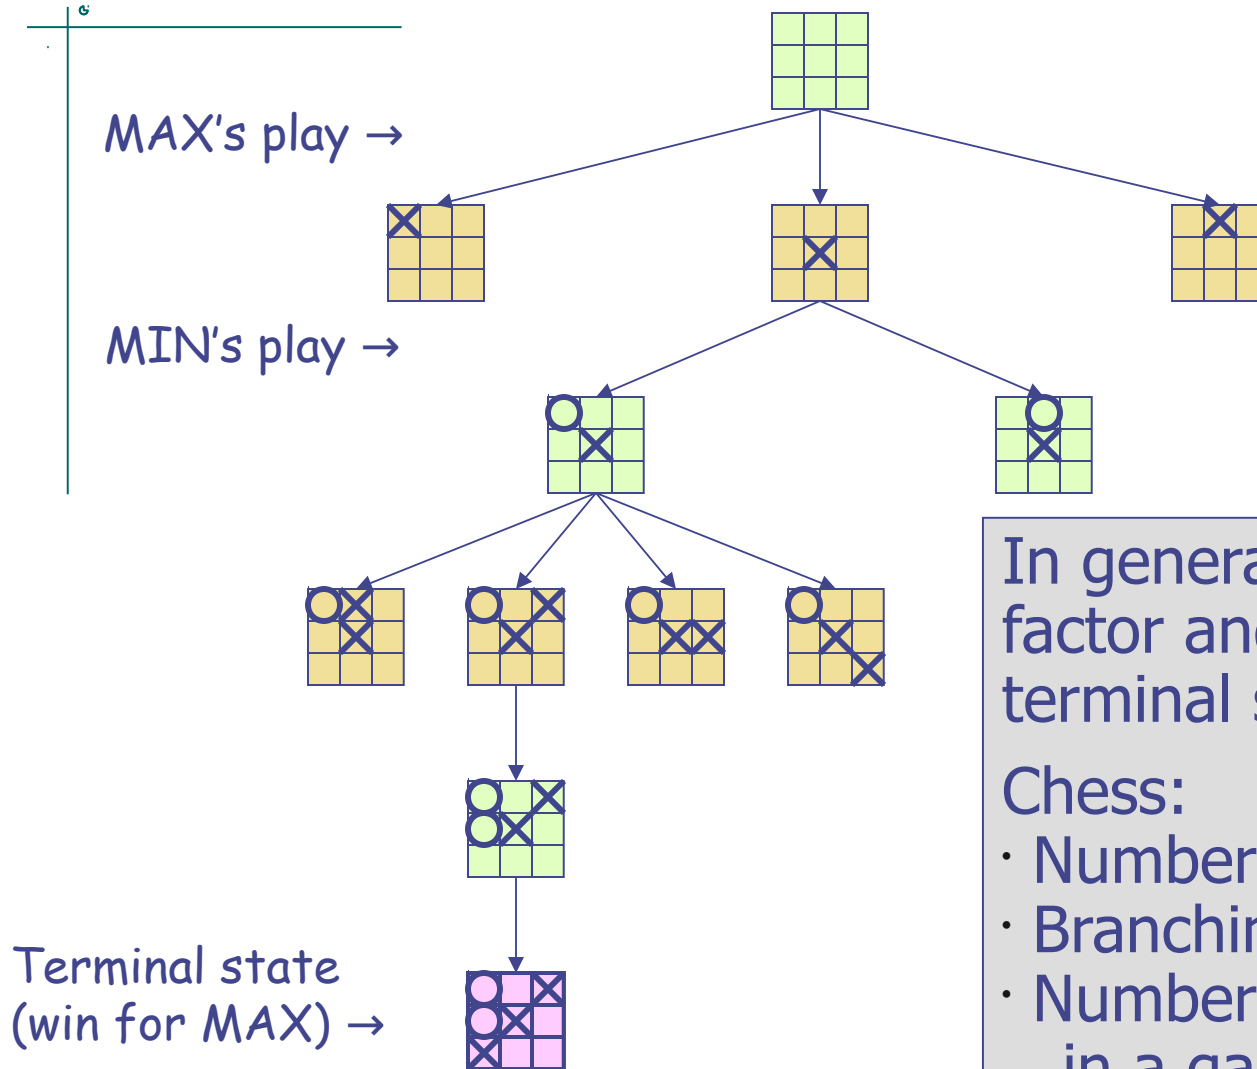
At each turn, the choice of which action to perform must be made within a **specified time limit**

The state space is enormous:  
only **a tiny fraction** of this space  
**can be explored** within the time limit

# Game Tree



# Game Tree



In general, the branching factor and the depth of terminal states are large

Chess:

- Number of states:  $\sim 10^{40}$
- Branching factor:  $\sim 35$
- Number of total moves in a game:  $\sim 100$



# Minimax Algorithm

- Expand the game tree uniformly from the current state (where it is MAX's turn to play) to depth  $h$  ("horizon")
- Compute evaluation function at every leaf
- Back-up the values from the leaves to the root:
  - MAX node → maximum evaluation of its successors
  - MIN node → minimum evaluation of its successors (assume the worst from MIN)
- Select the move toward a MIN node that has the largest backed-up value

# Evaluation Function

- ◆ Function **e**: State  $\rightarrow$  Number
- ◆  $e(s)$  estimates how favorable  $s$  is for MAX


$e(s) > 0$  means that  $s$  is favorable to MAX  
(the larger the better)

$e(s) < 0$  means that  $s$  is favorable to MIN


$e(s) = 0$  means that  $s$  is neutral

# Example: Tic-tac-Toe


$e(s)$  = number of rows, columns,  
and diagonals open for MAX  
- number of rows, columns,  
and diagonals open for MIN



$$8 - 8 = 0$$



$$6 - 4 = 2$$



$$3 - 3 = 0$$

# Creating an Evaluation Function

Usually a weighted sum of “features”:


$$e(s) = \sum_{i=1}^n w_i * f_i(s)$$

Features may include:

- Number of pieces of each type
- Number of possible moves
- Number of squares controlled

# Backing up Values

Tic-Tac-Toe tree  
at horizon = 2



# Why using backed-up values?

At each non-leaf node  $N$ , the backed-up value is the value of the **best state that MAX can reach** at depth  $h$  if MIN plays well (by the same criterion as MAX applies to itself)

If  $e$  is to be trusted in the first place, then the backed-up value is a better estimate of how favorable  $\text{STATE}(N)$  is than  $e(\text{STATE}(N))$

# Game Playing (for MAX)

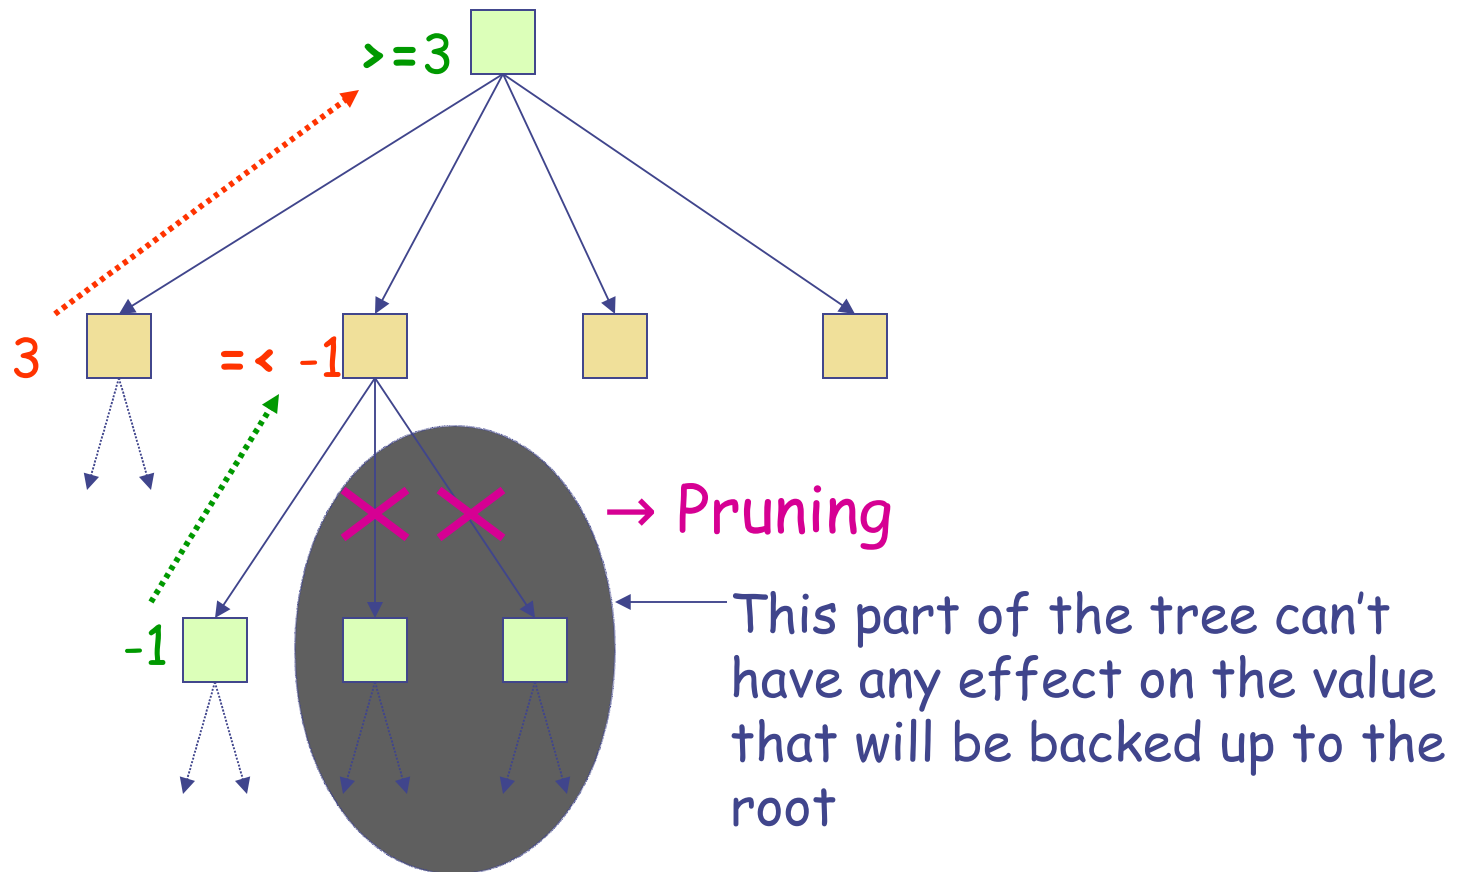
- Repeat until a terminal state is reached
  - Select move using Minimax
  - Execute move
  - Observe MIN's move

Note that at each cycle the large game tree built to horizon  $h$  is used to select only one move

All is repeated again at the next cycle (a sub-tree of depth  $h-2$  can be re-used)

# Can we do better?


Yes ! Much better !






# Example

6




# Example

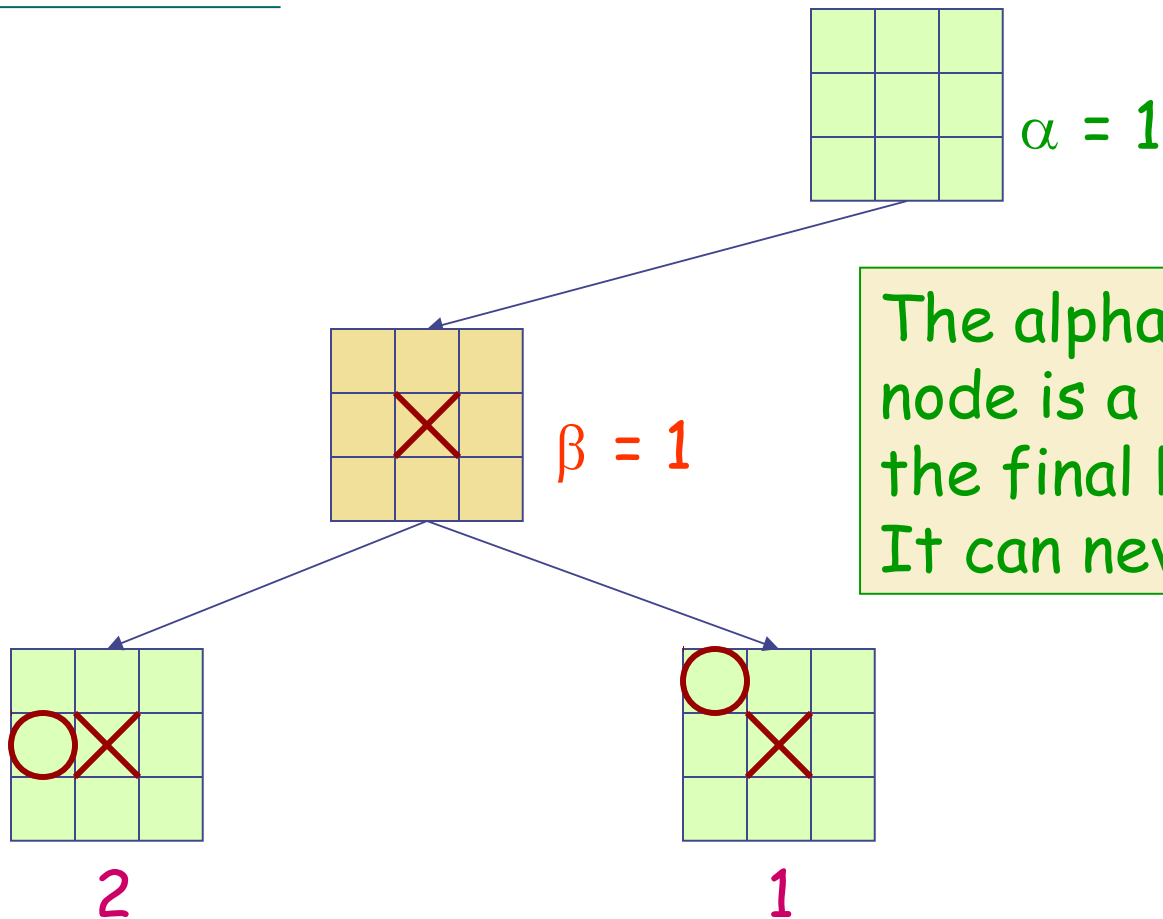


The beta value of a MIN node is an upper bound on the final backed-up value. It can never increase


# Example




# Example



# Example



# Example



# Alpha-Beta Pruning

Explore the game tree to **depth h** in **depth-first** manner

**Back up** alpha and beta values whenever possible

**Prune branches** that can't lead to changing the final decision

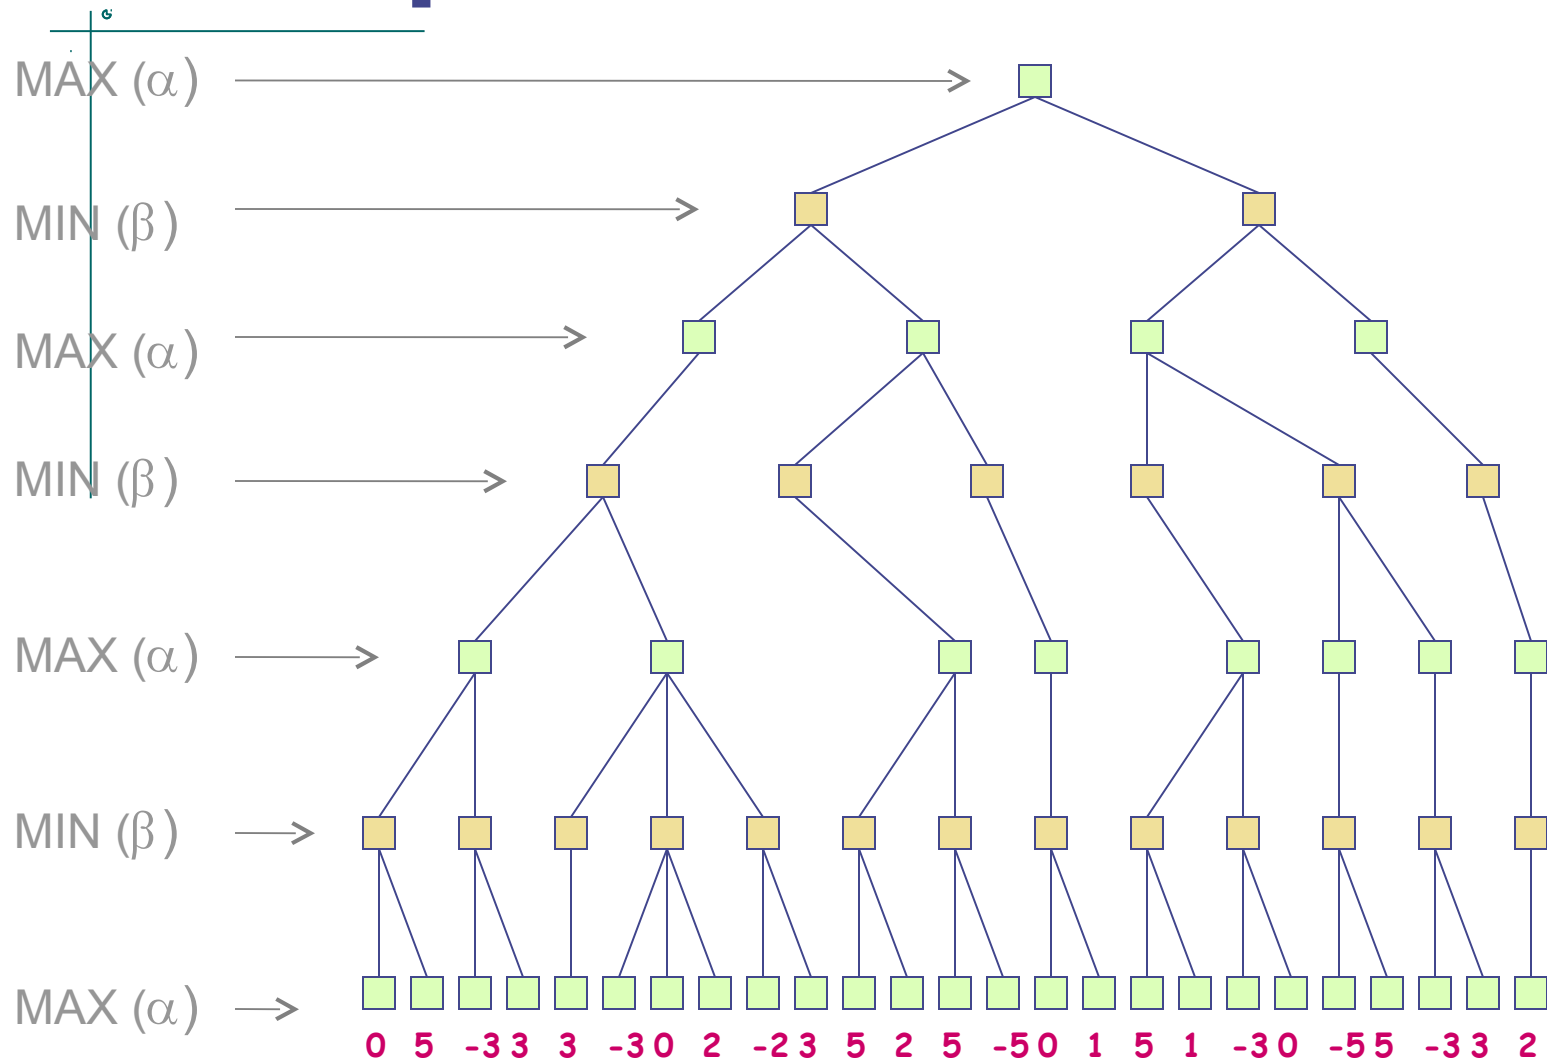
# Alpha-Beta Algorithm

Update the alpha/beta value of the parent of a node N when the search below N has been completed or discontinued

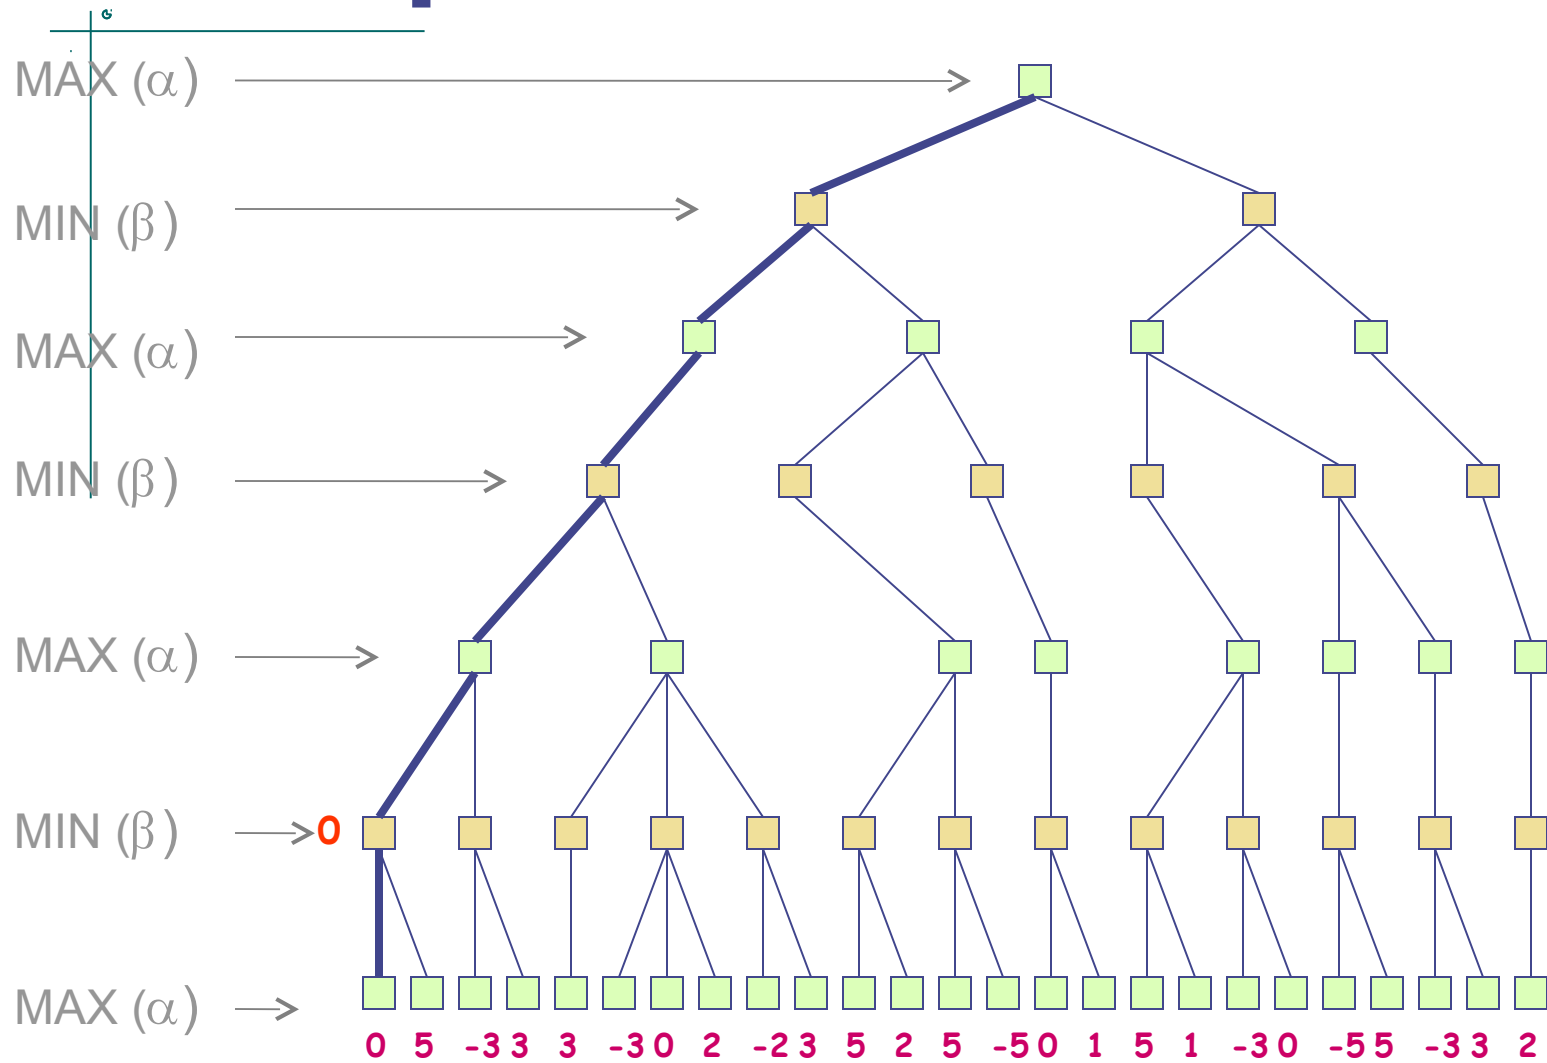
- Discontinue the search below a **MAX** node N if its **alpha value** is  $\geq$  the **beta value** of a MIN ancestor of N
- Discontinue the search below a **MIN** node N if its **beta value** is  $\leq$  the **alpha value** of a MAX ancestor of N



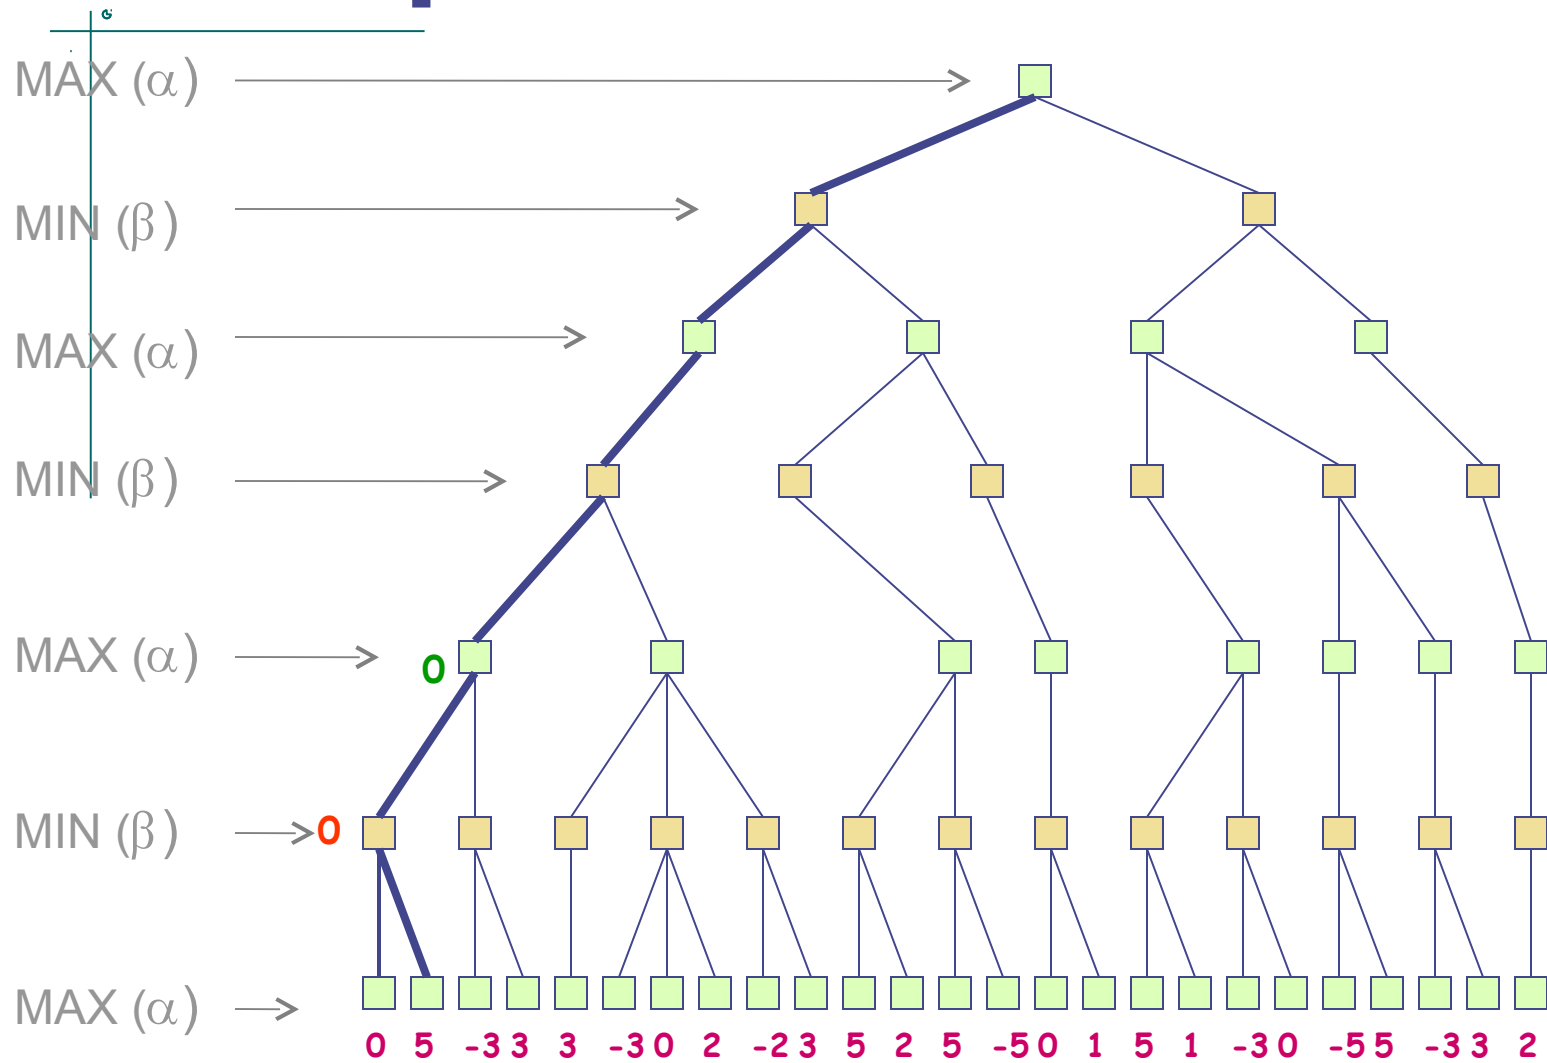
# Example



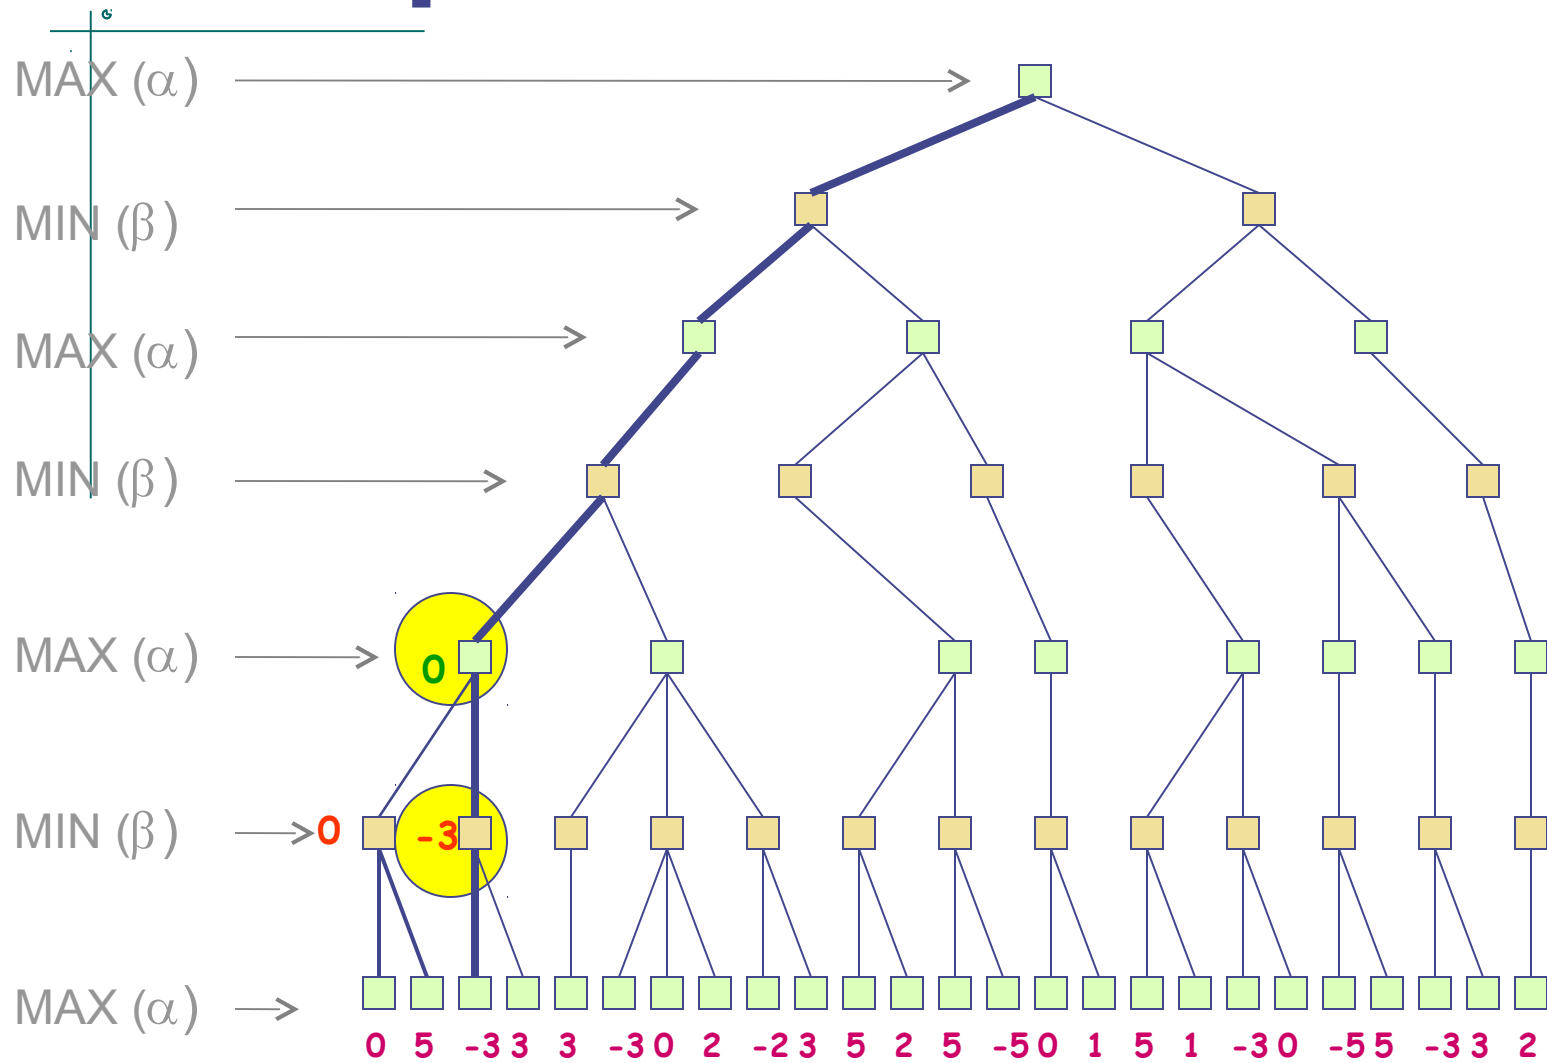
# Example



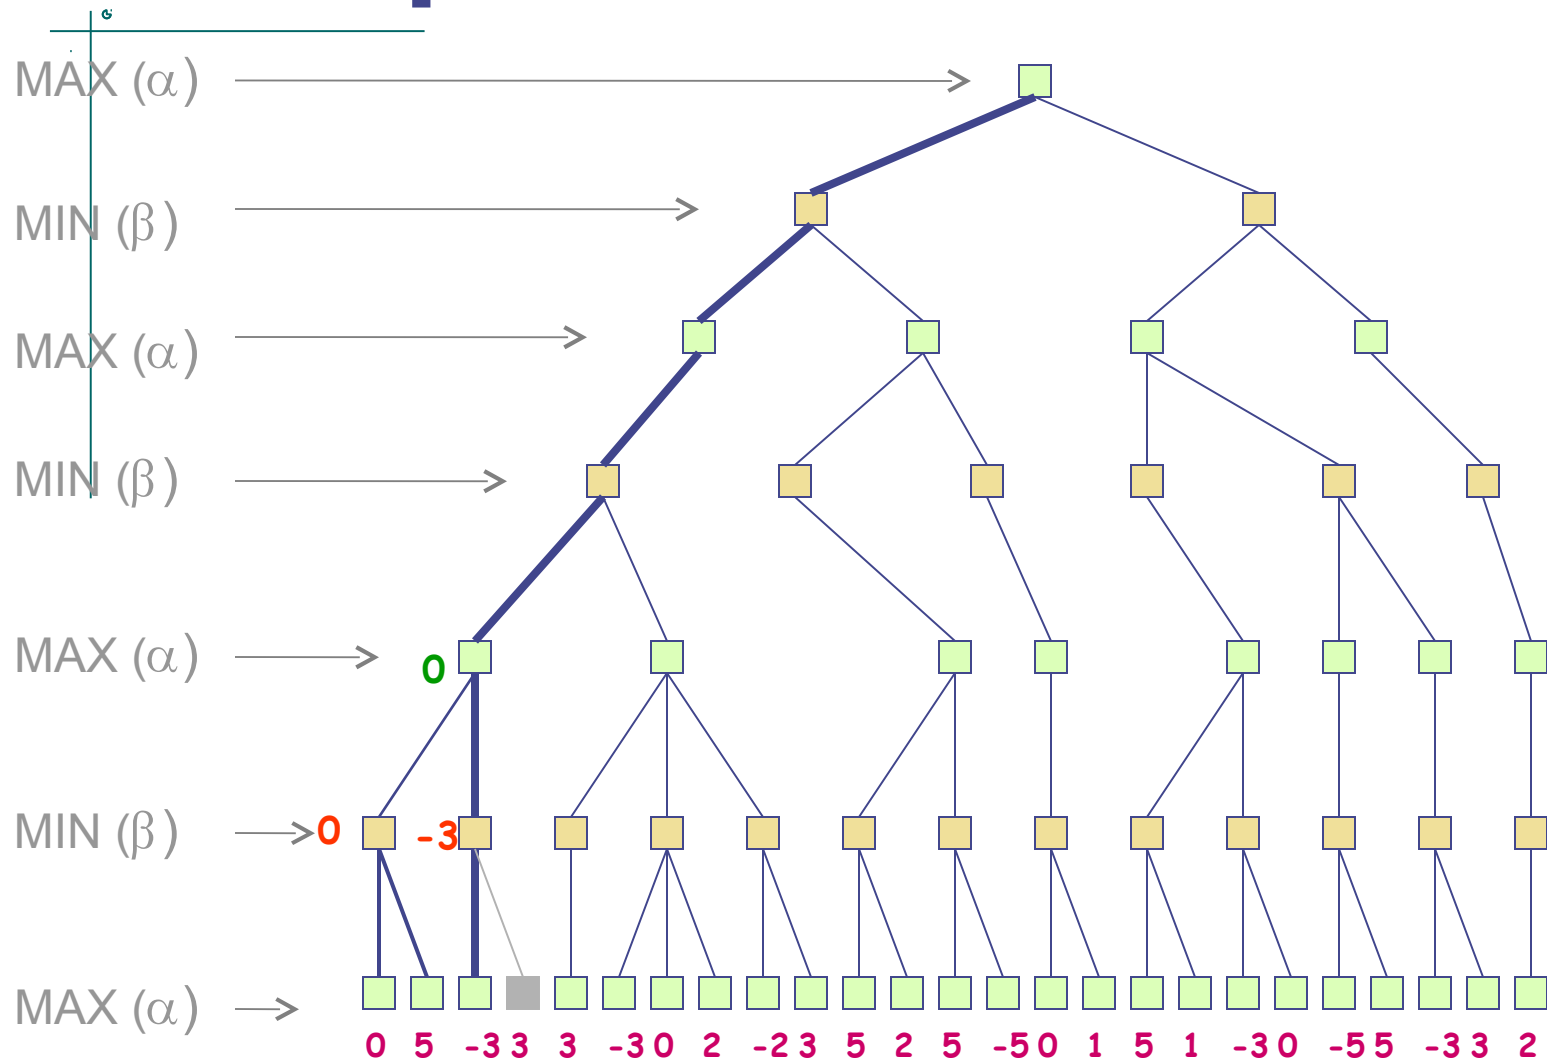
# Example



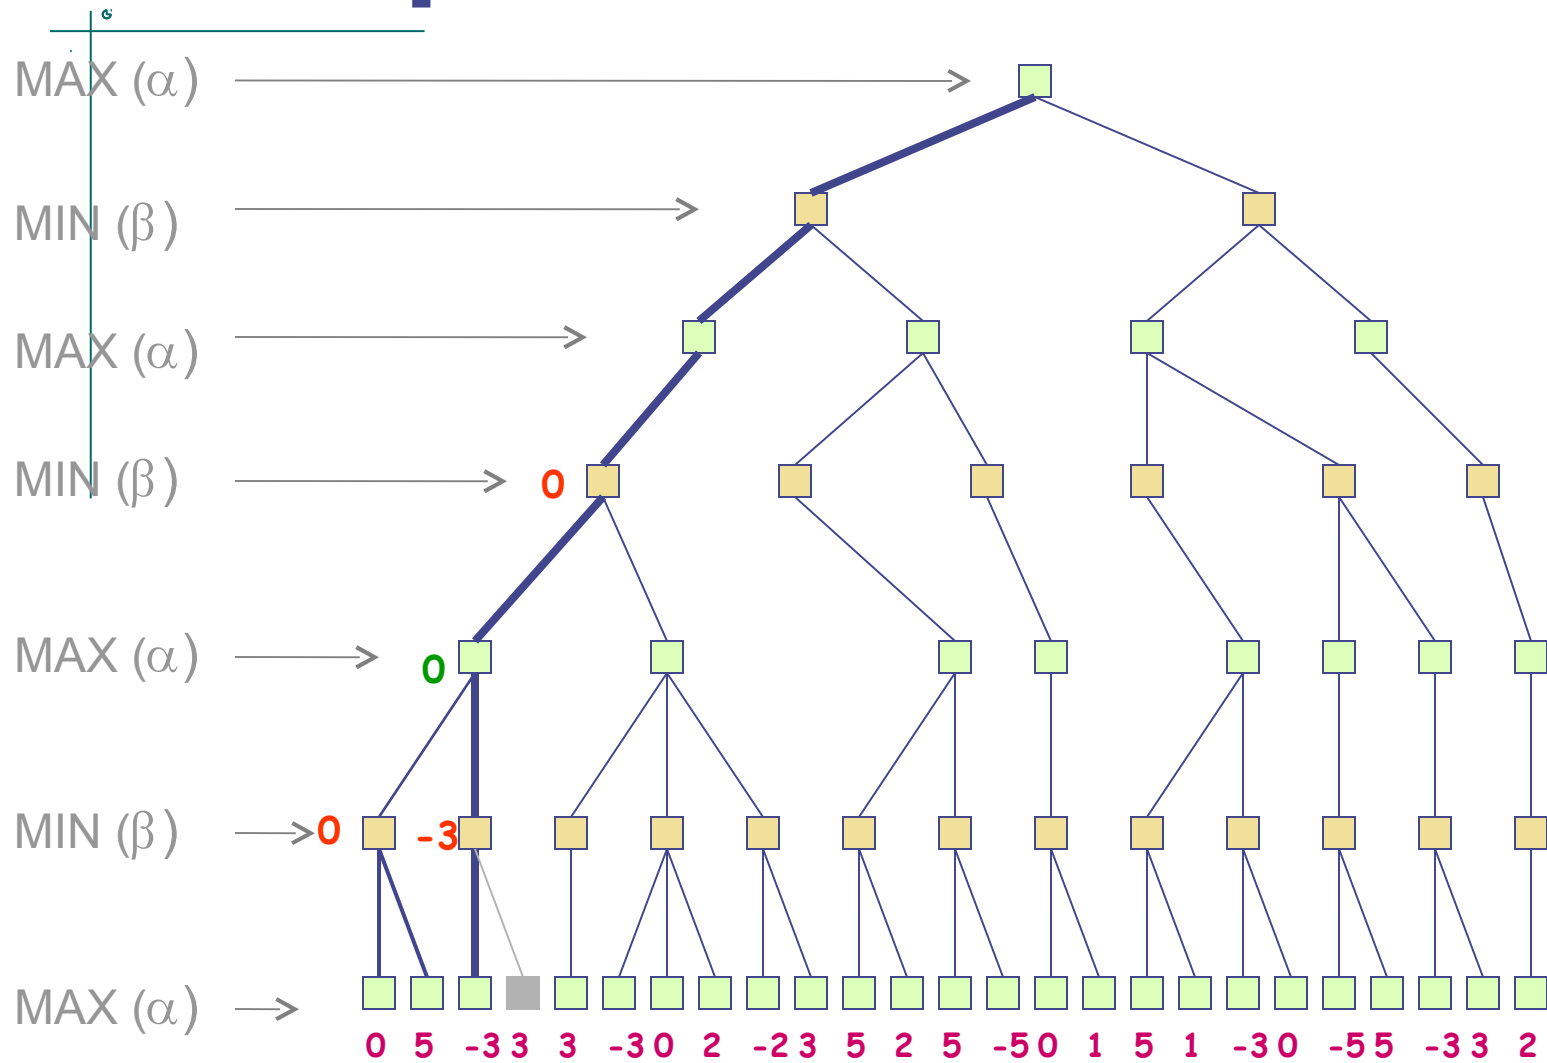
# Example



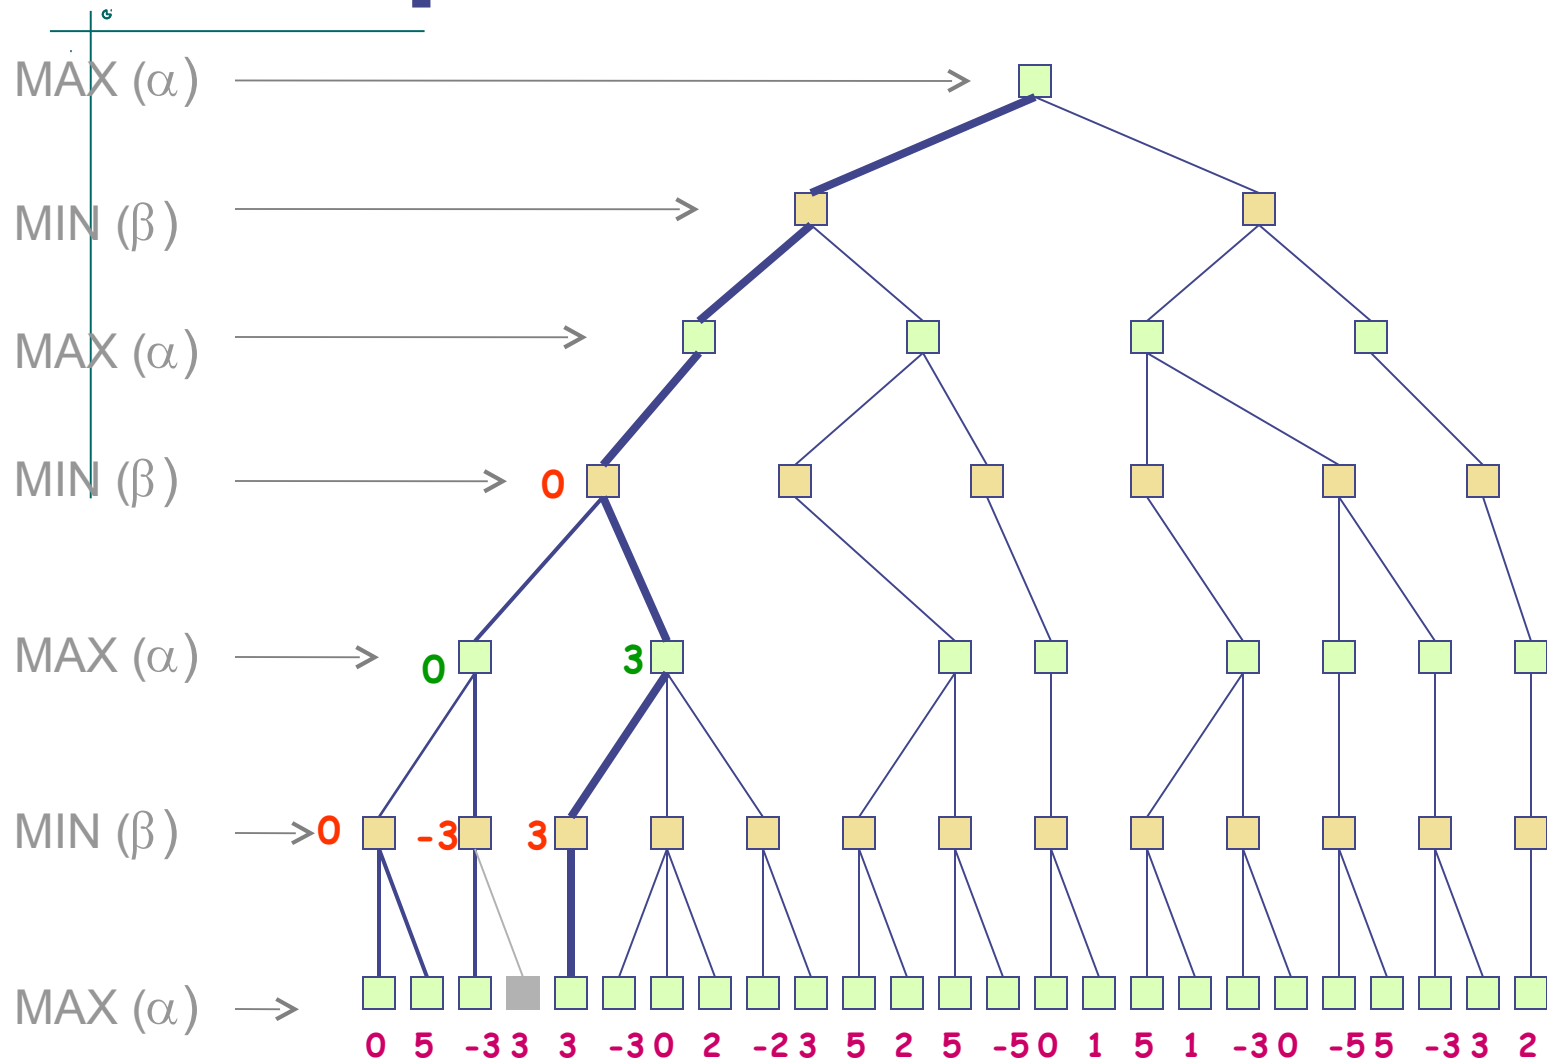
# Example



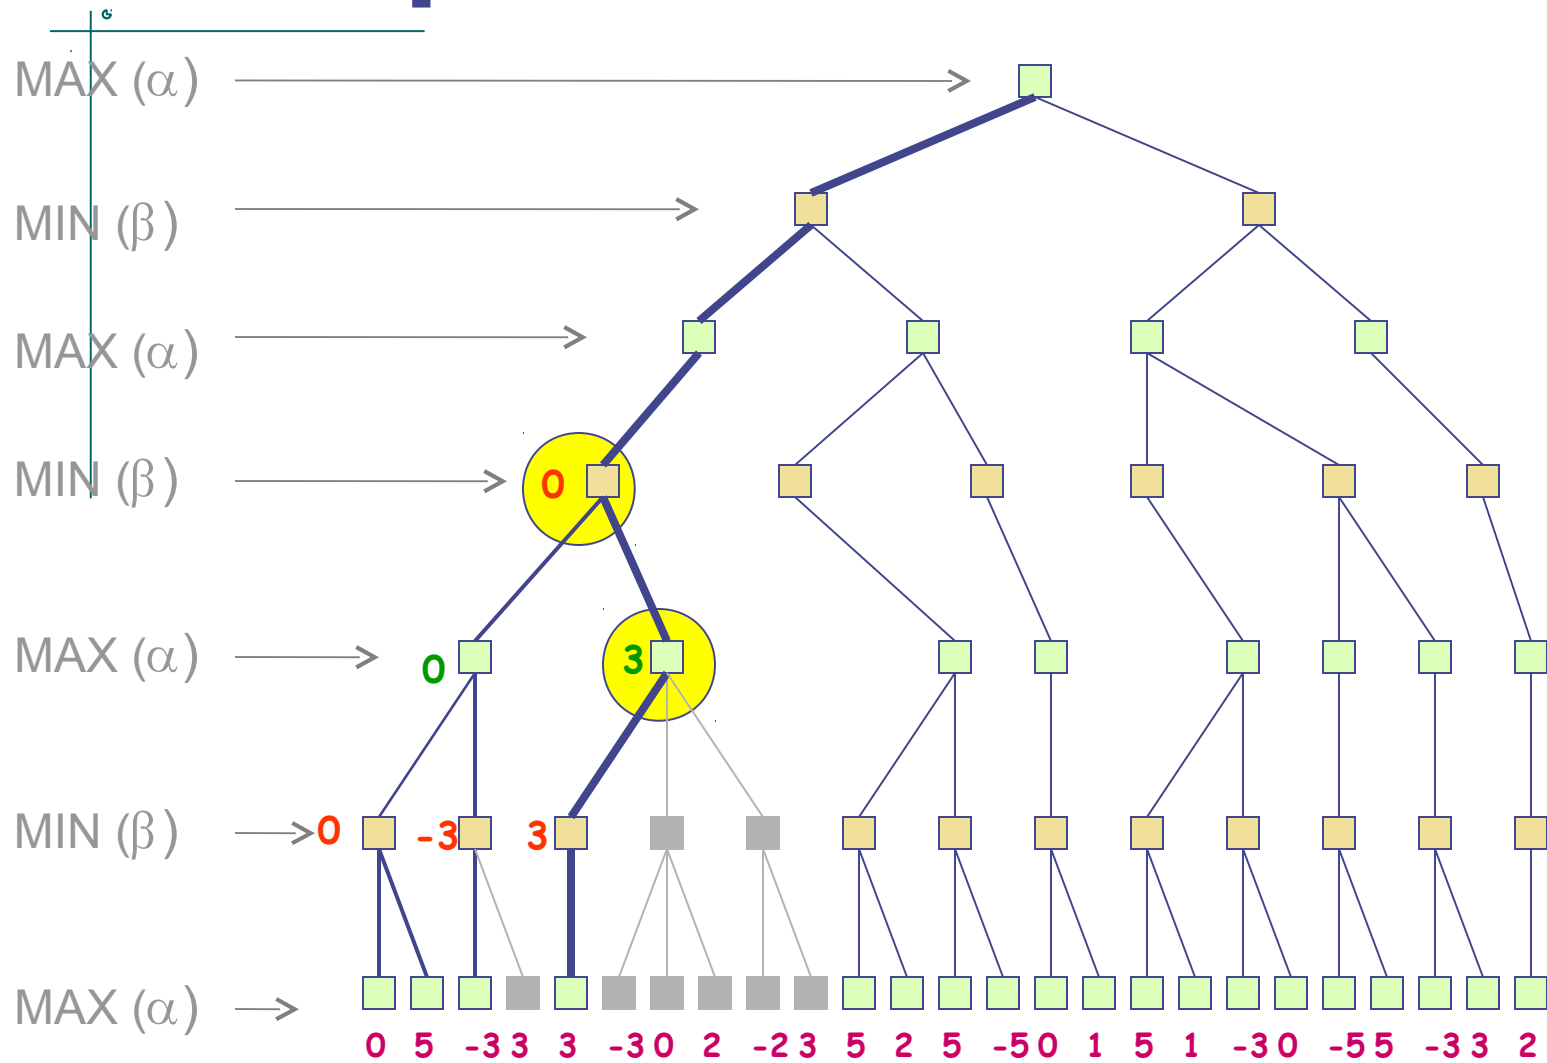
# Example



# Example

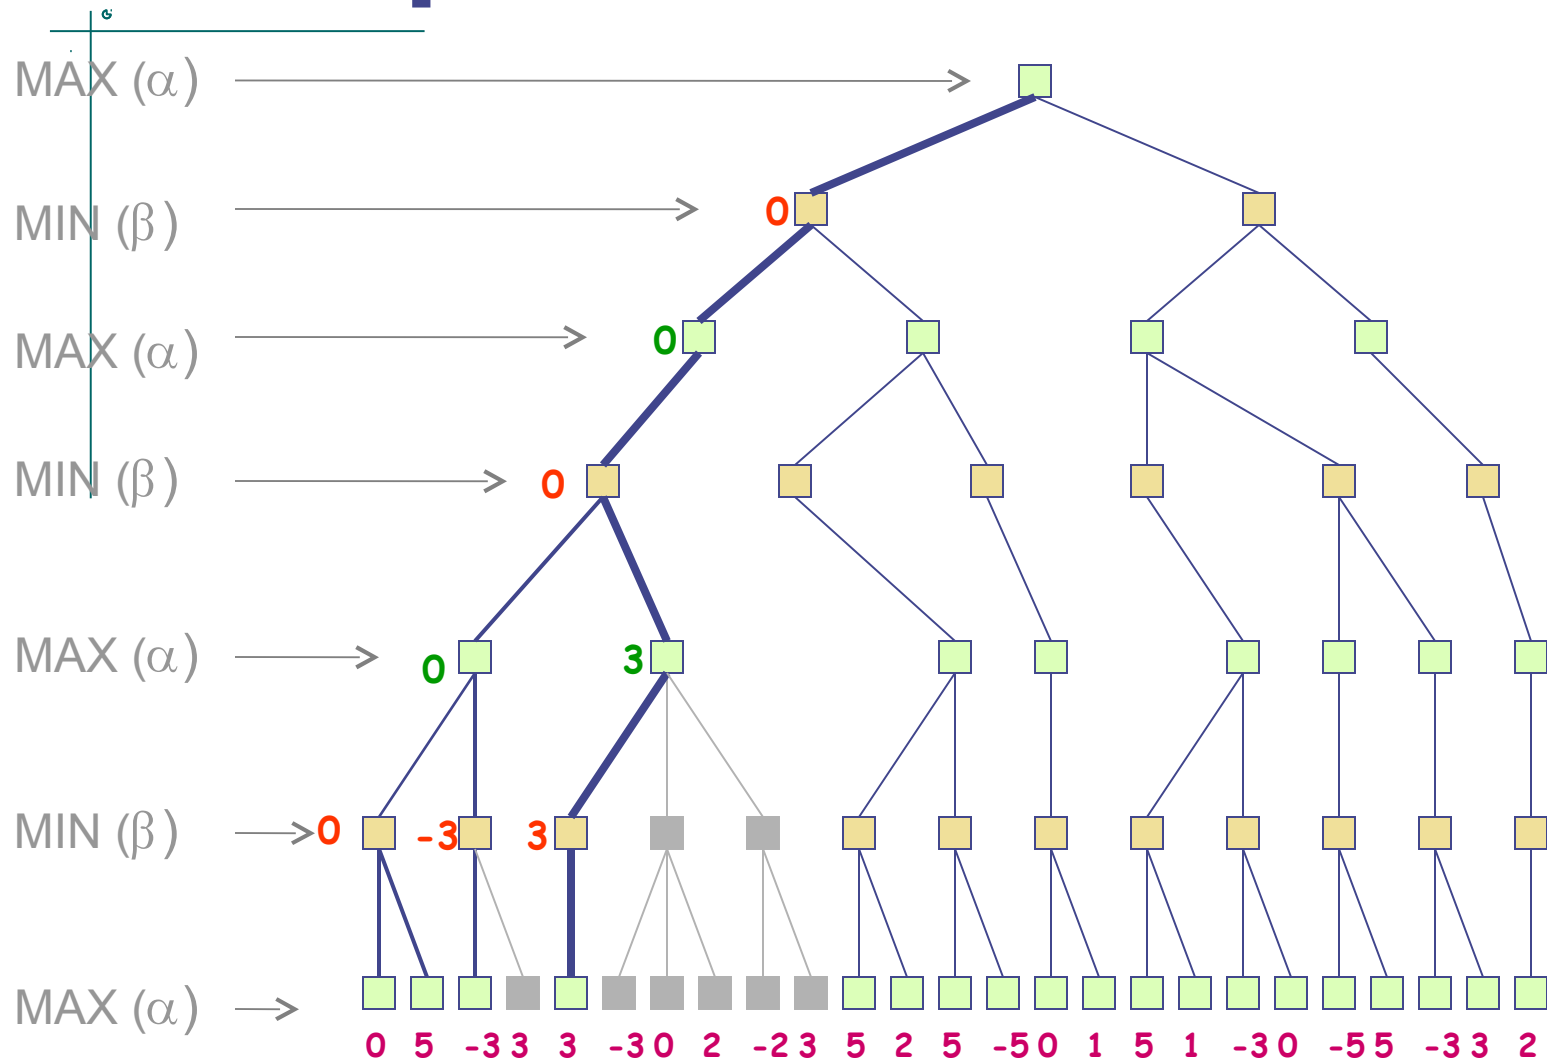


# Example

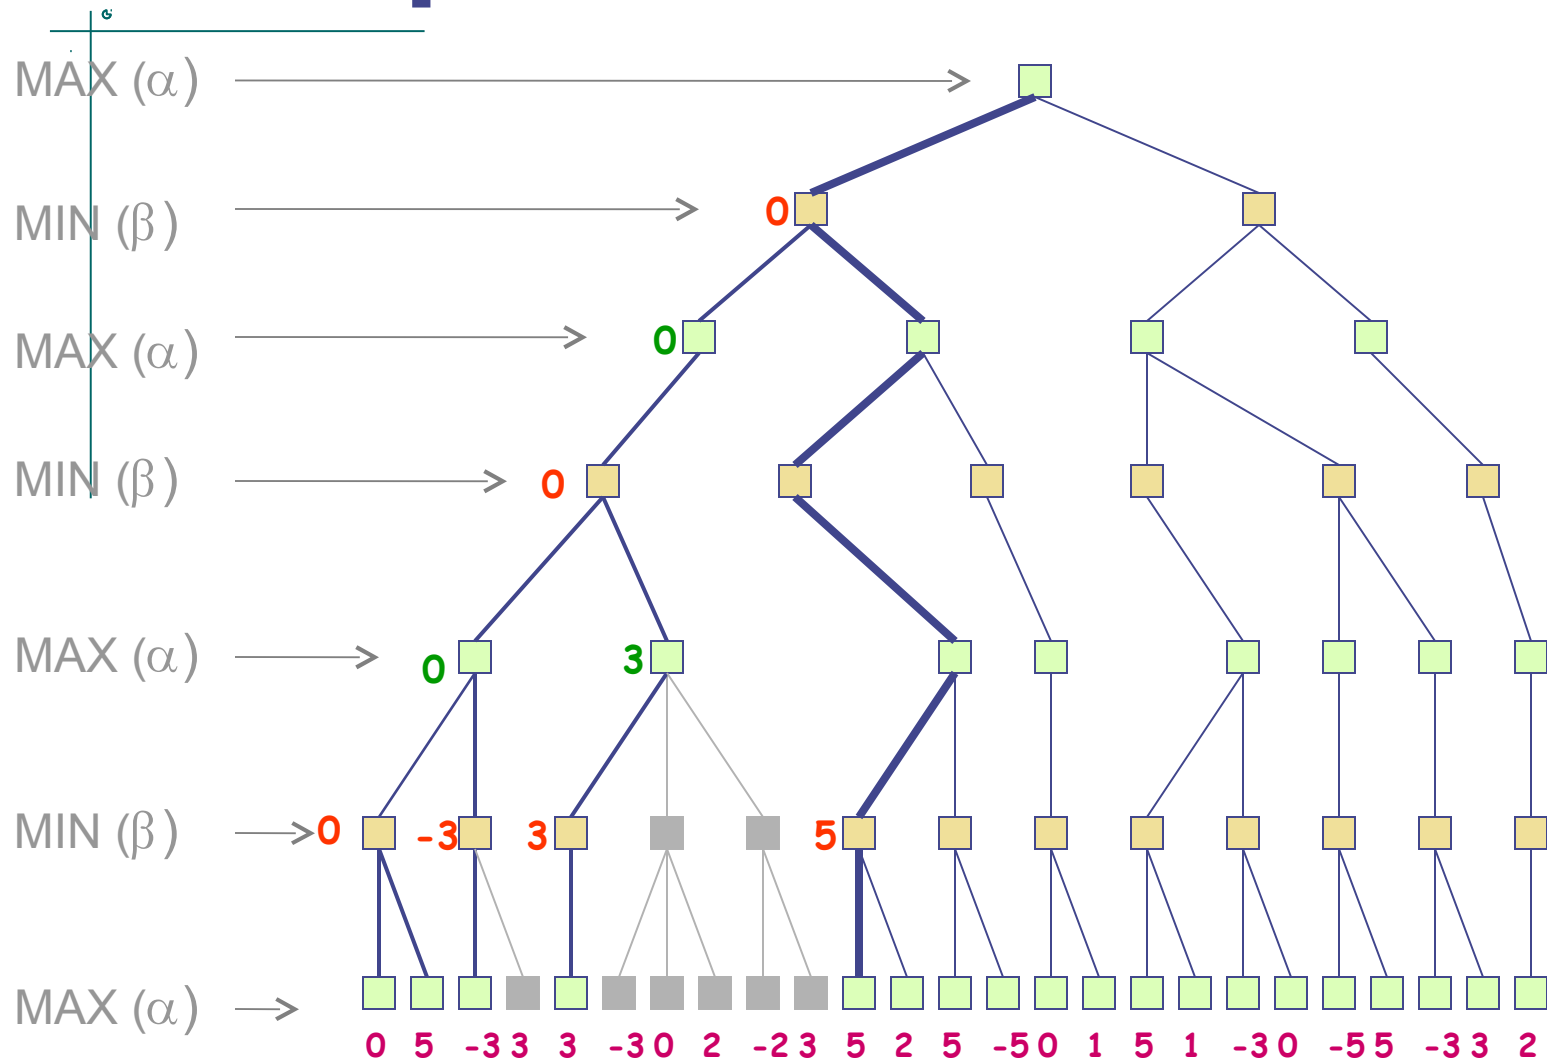




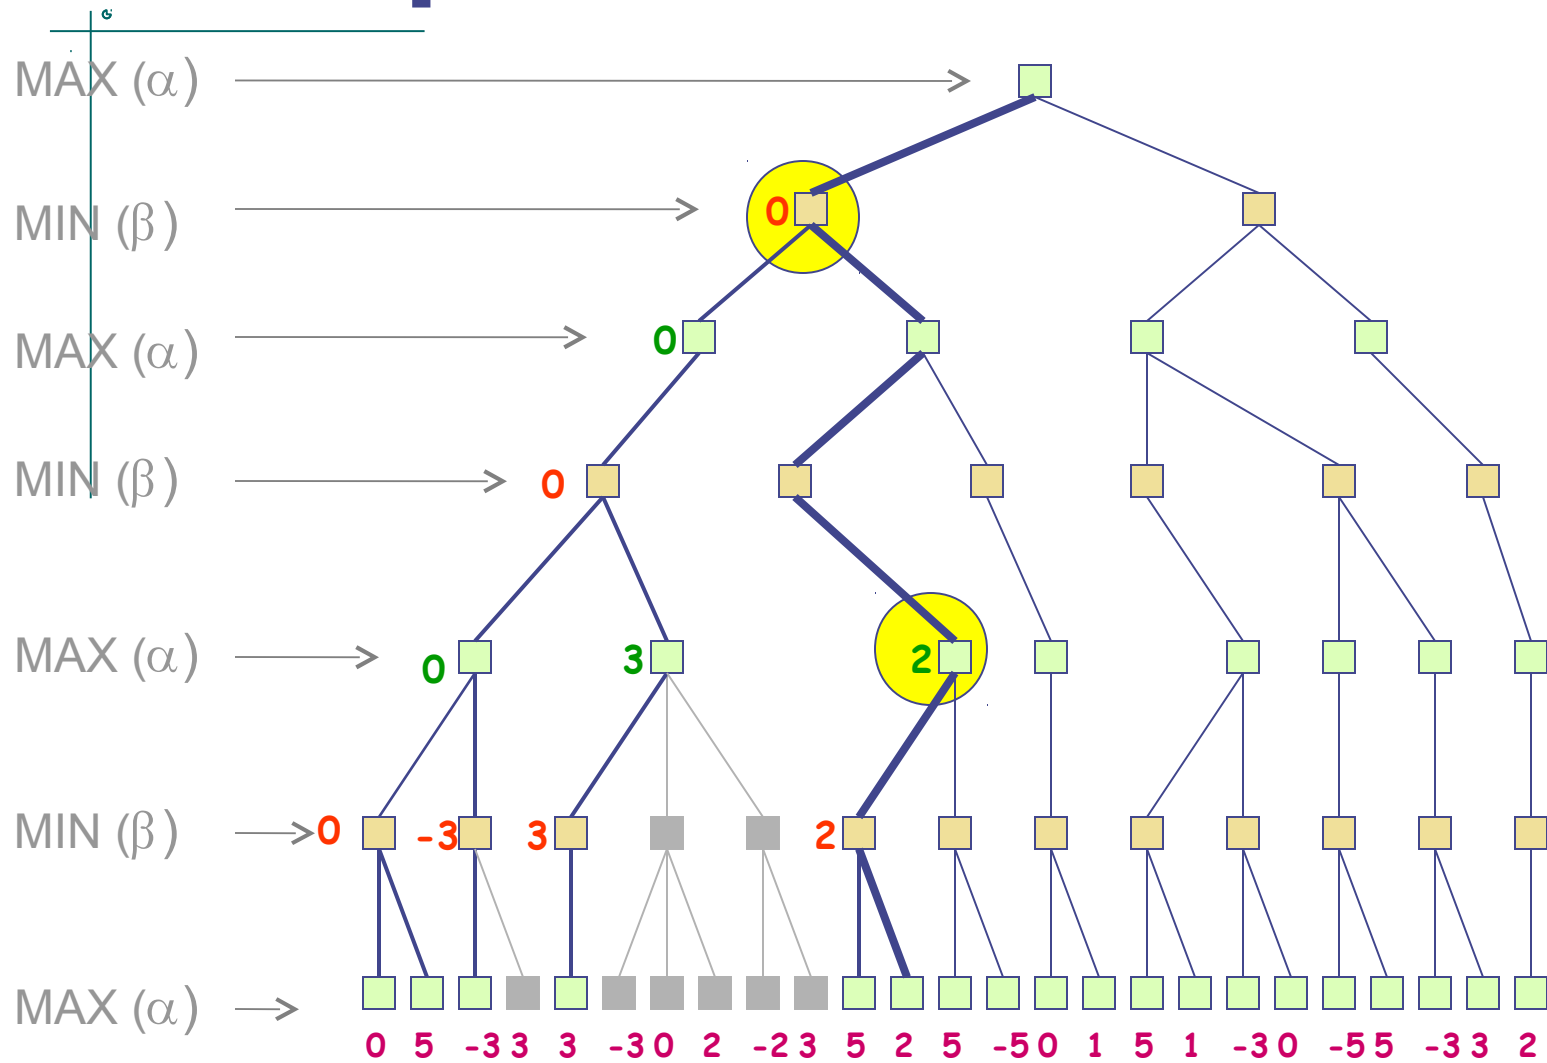
# Example



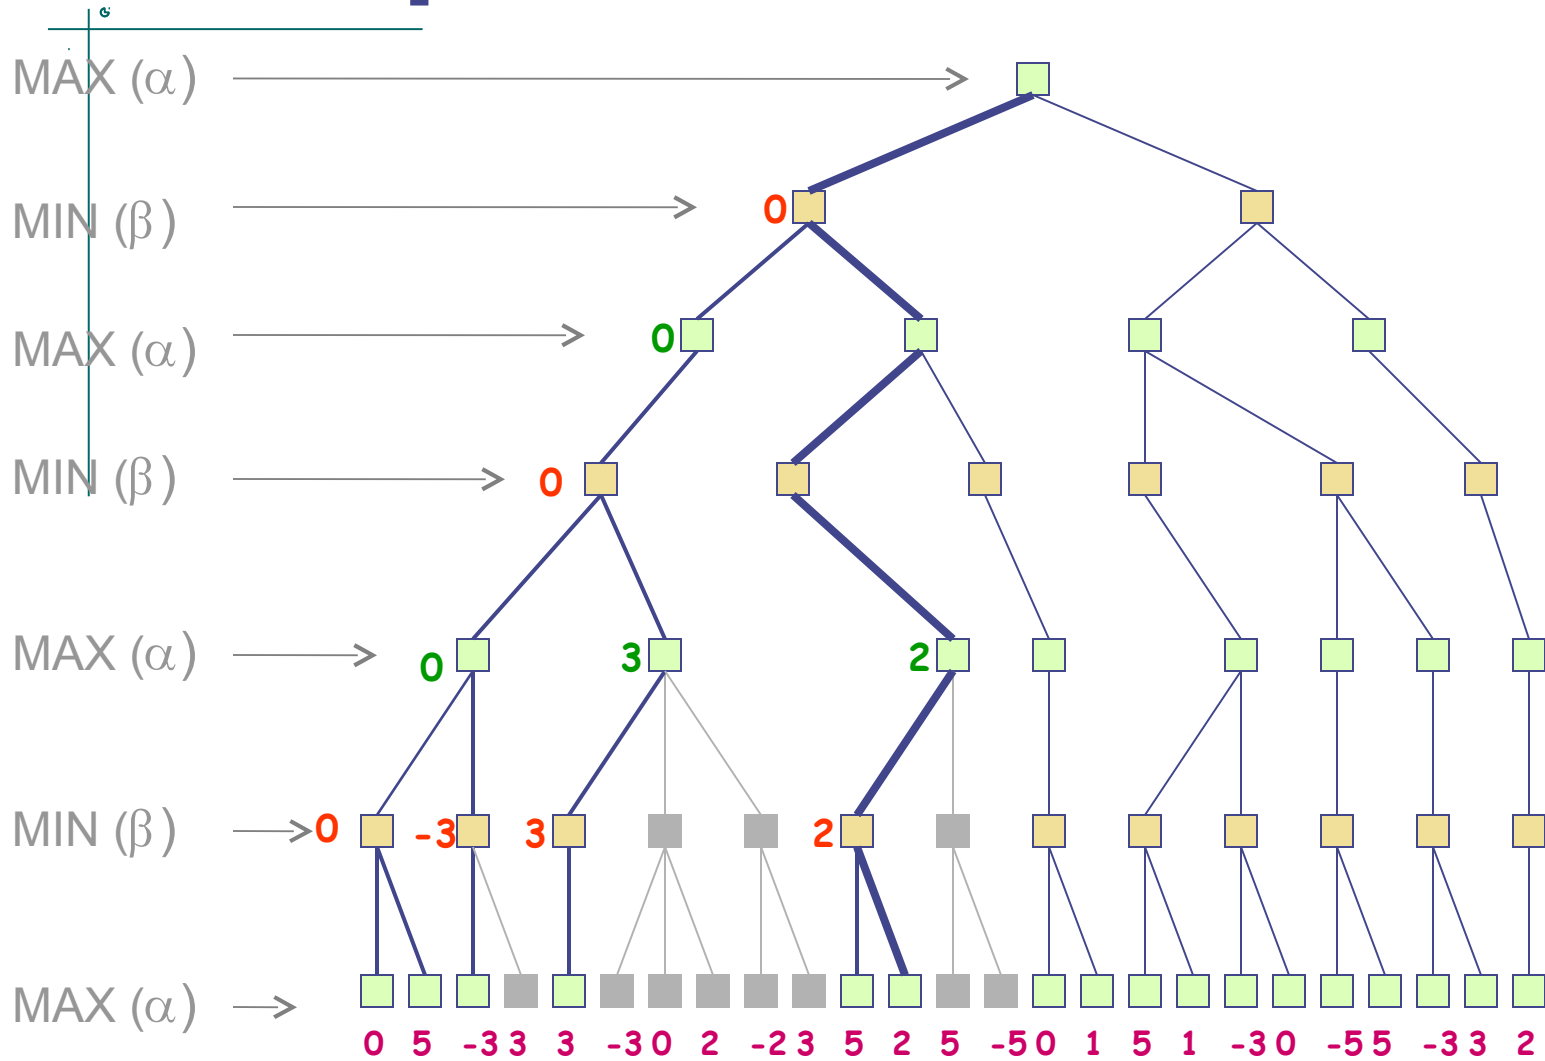
# Example



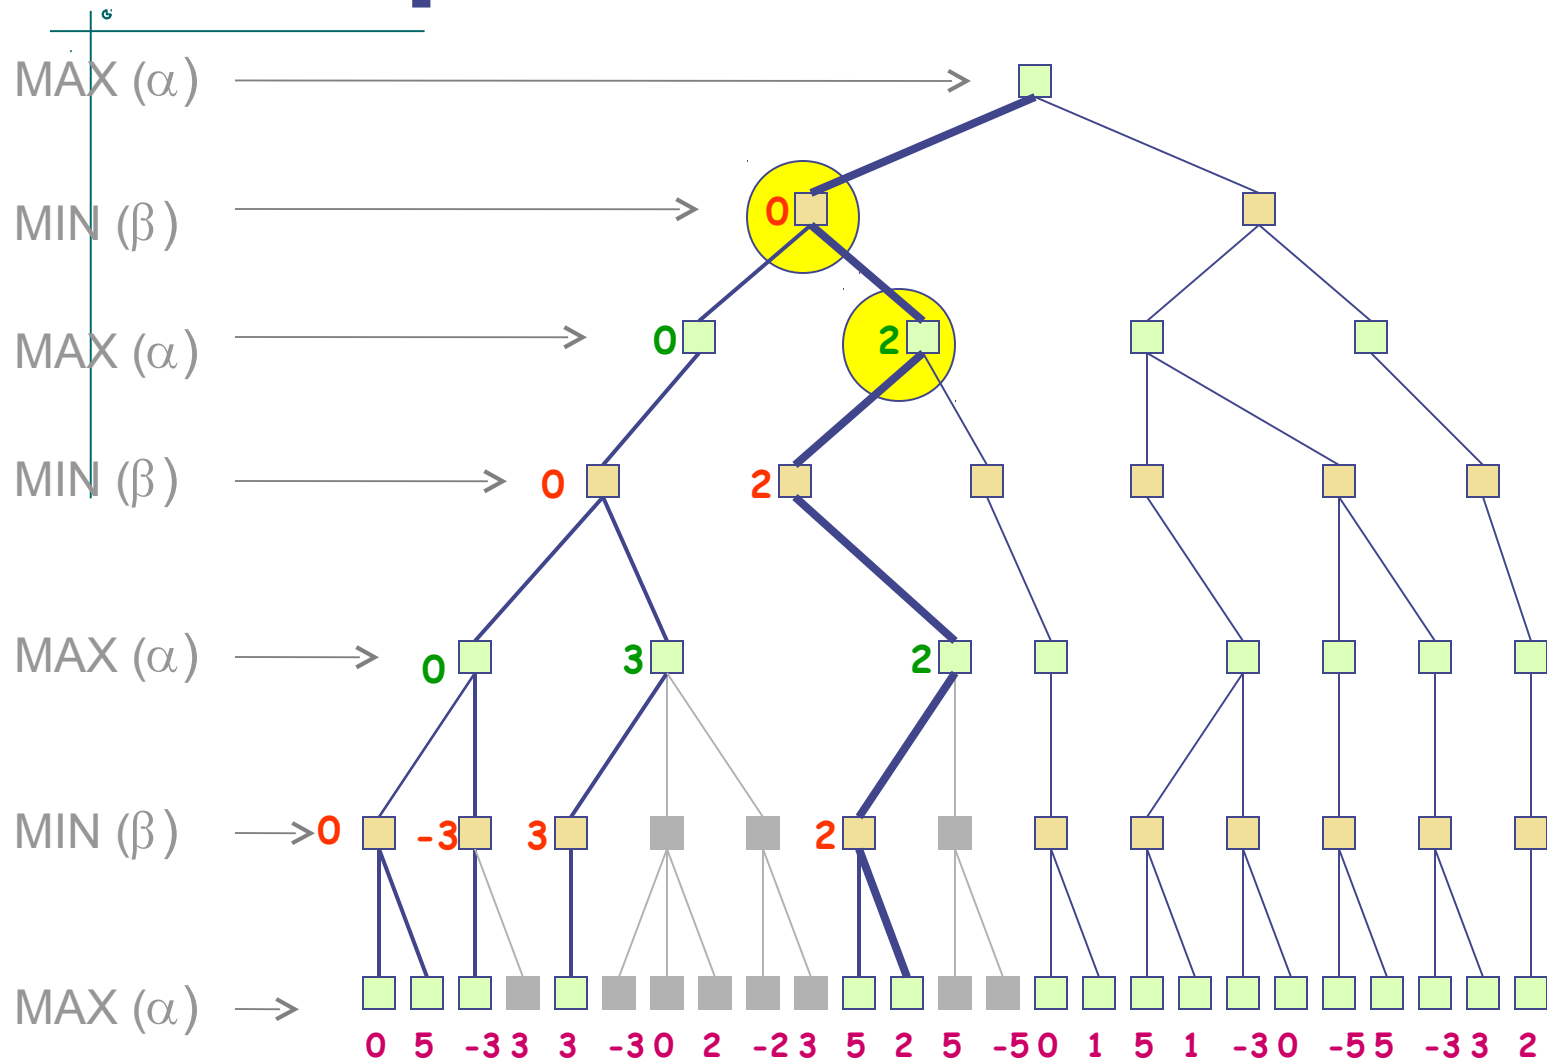
# Example



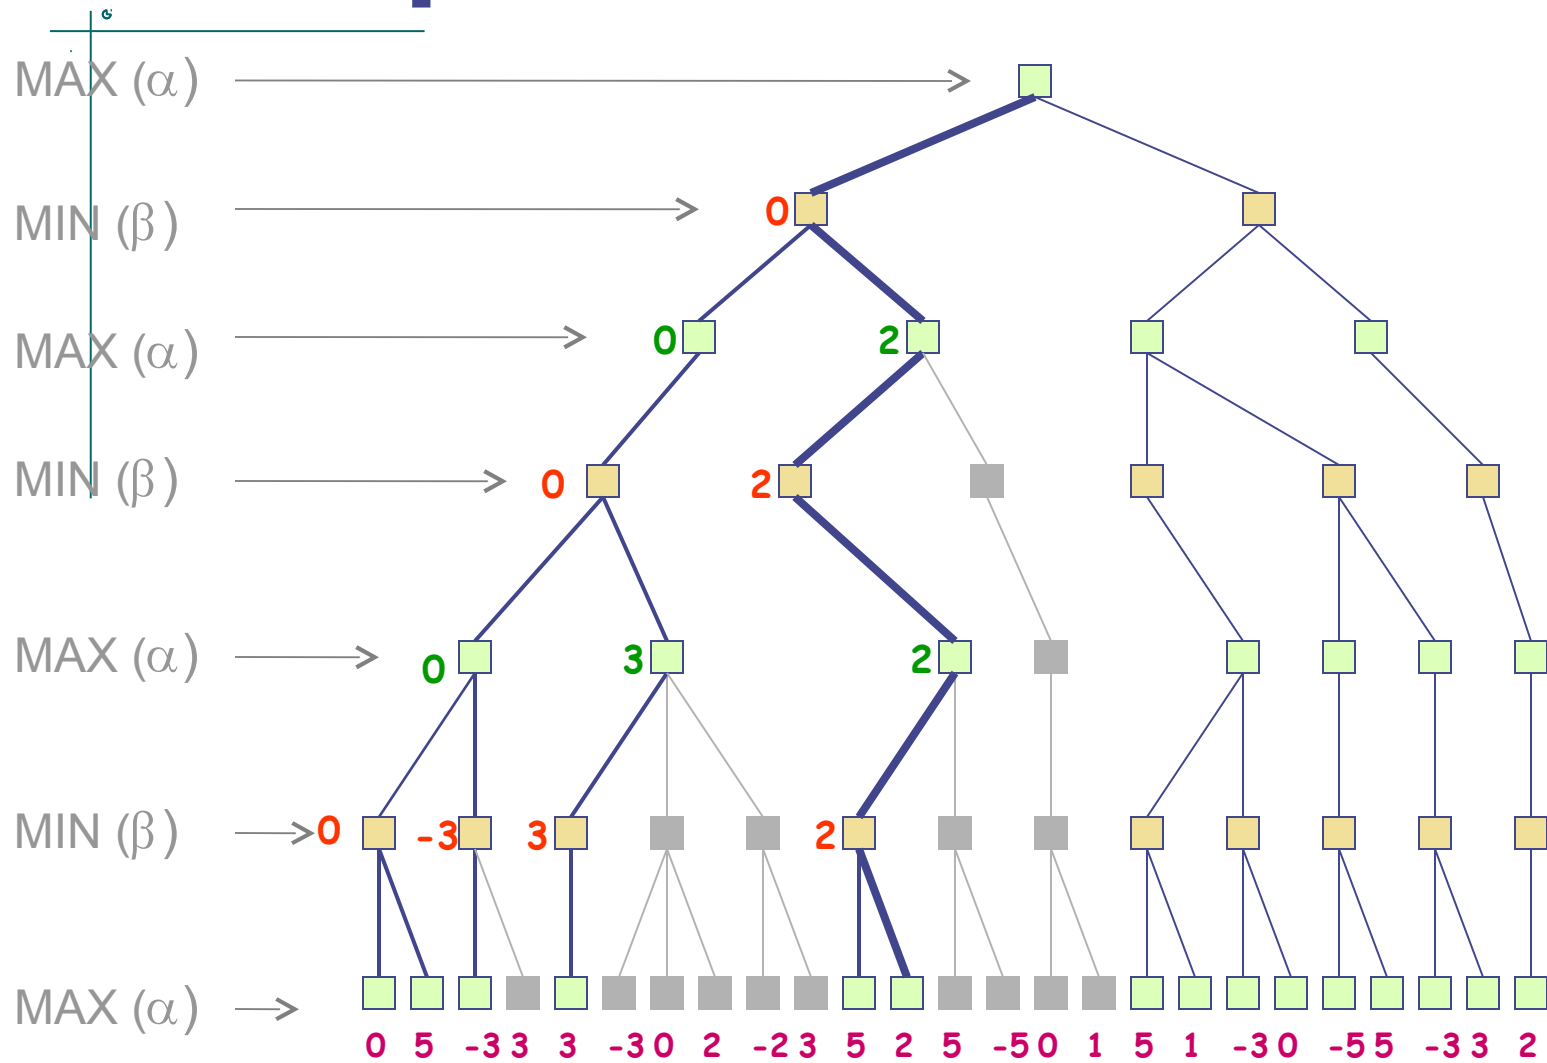
# Example



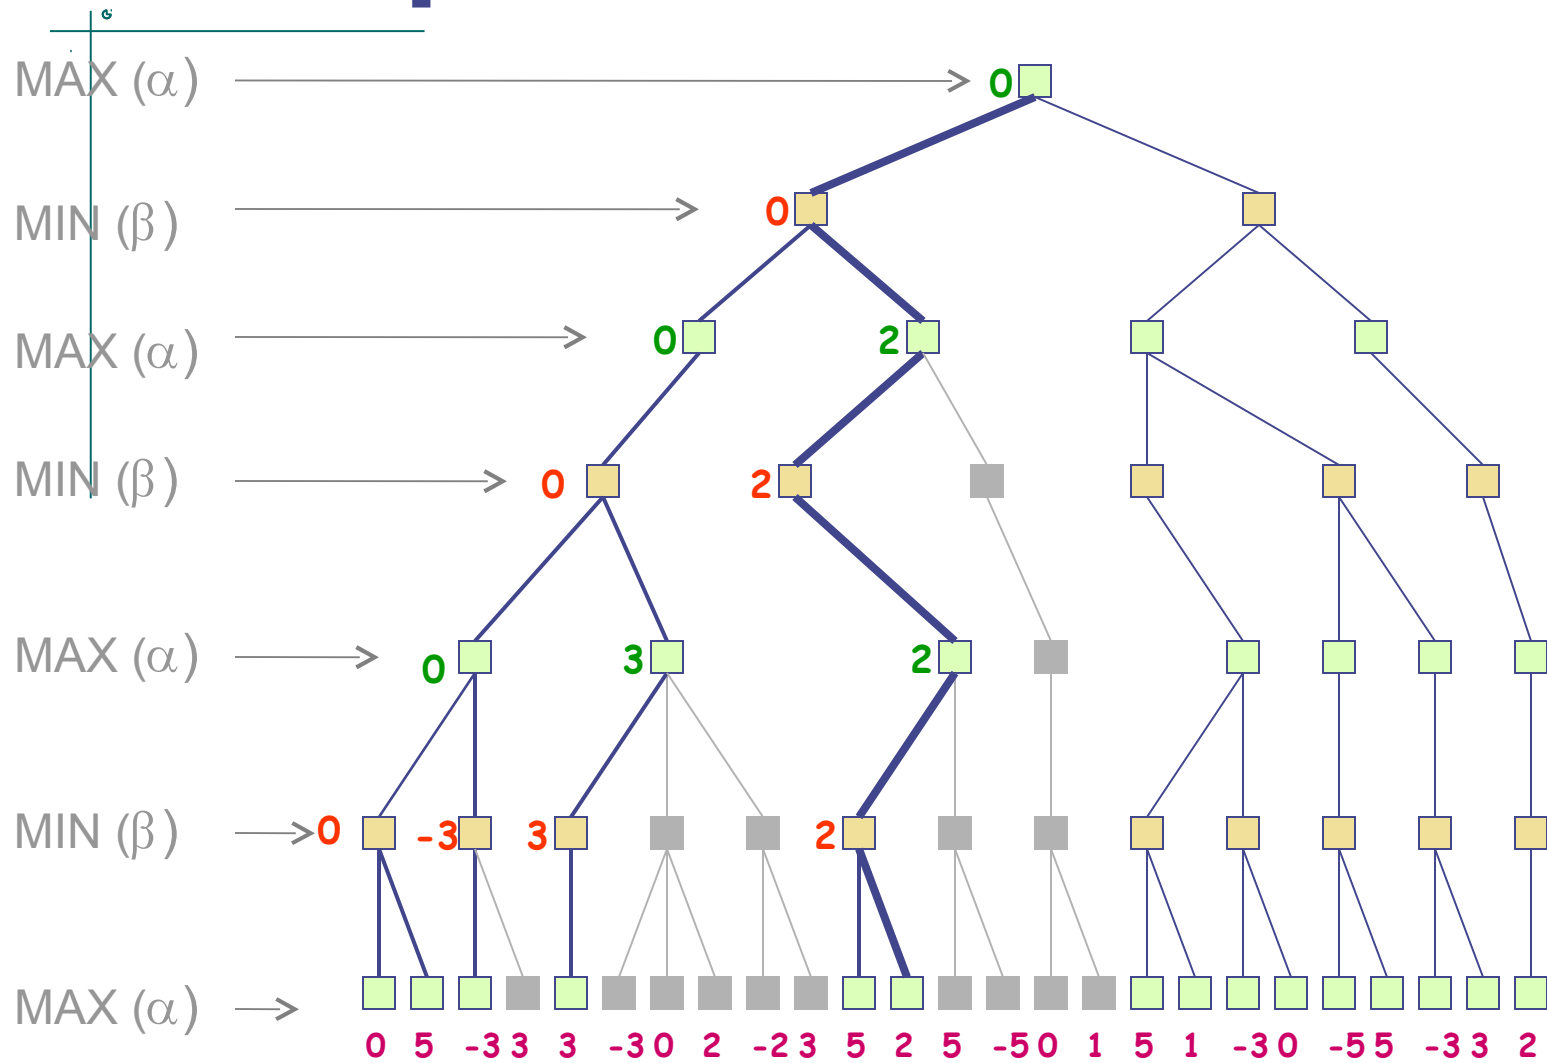
# Example




# Example



# Example




# Example







# Example




# Example




3



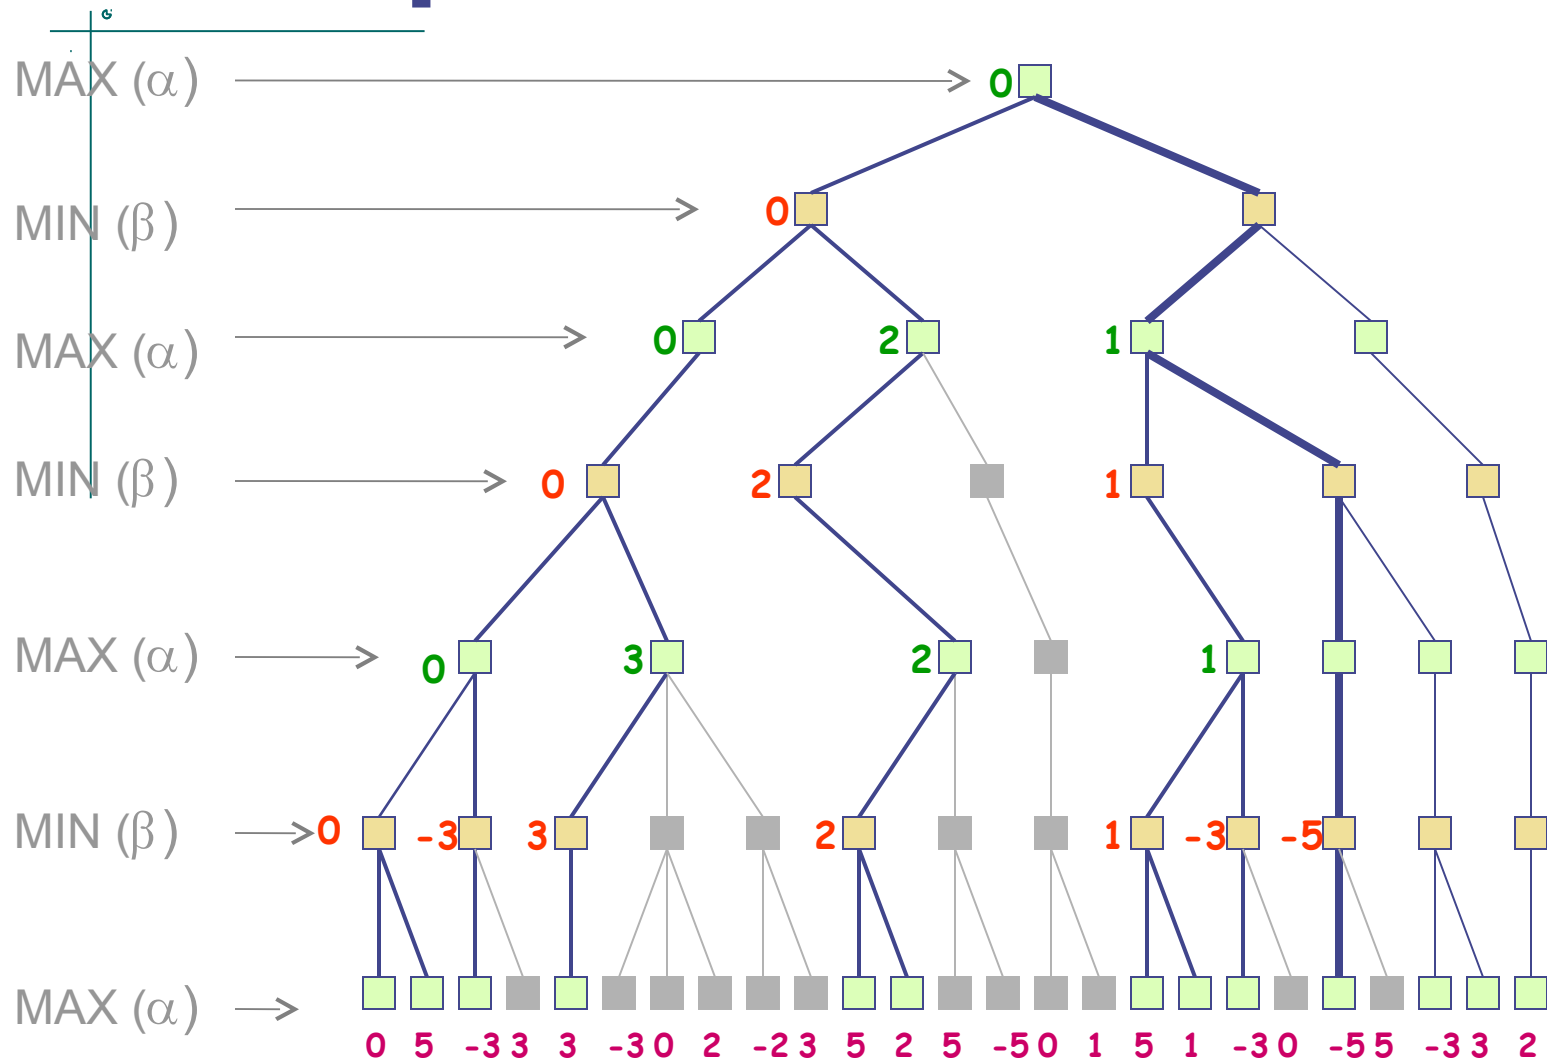
# Example




# Example



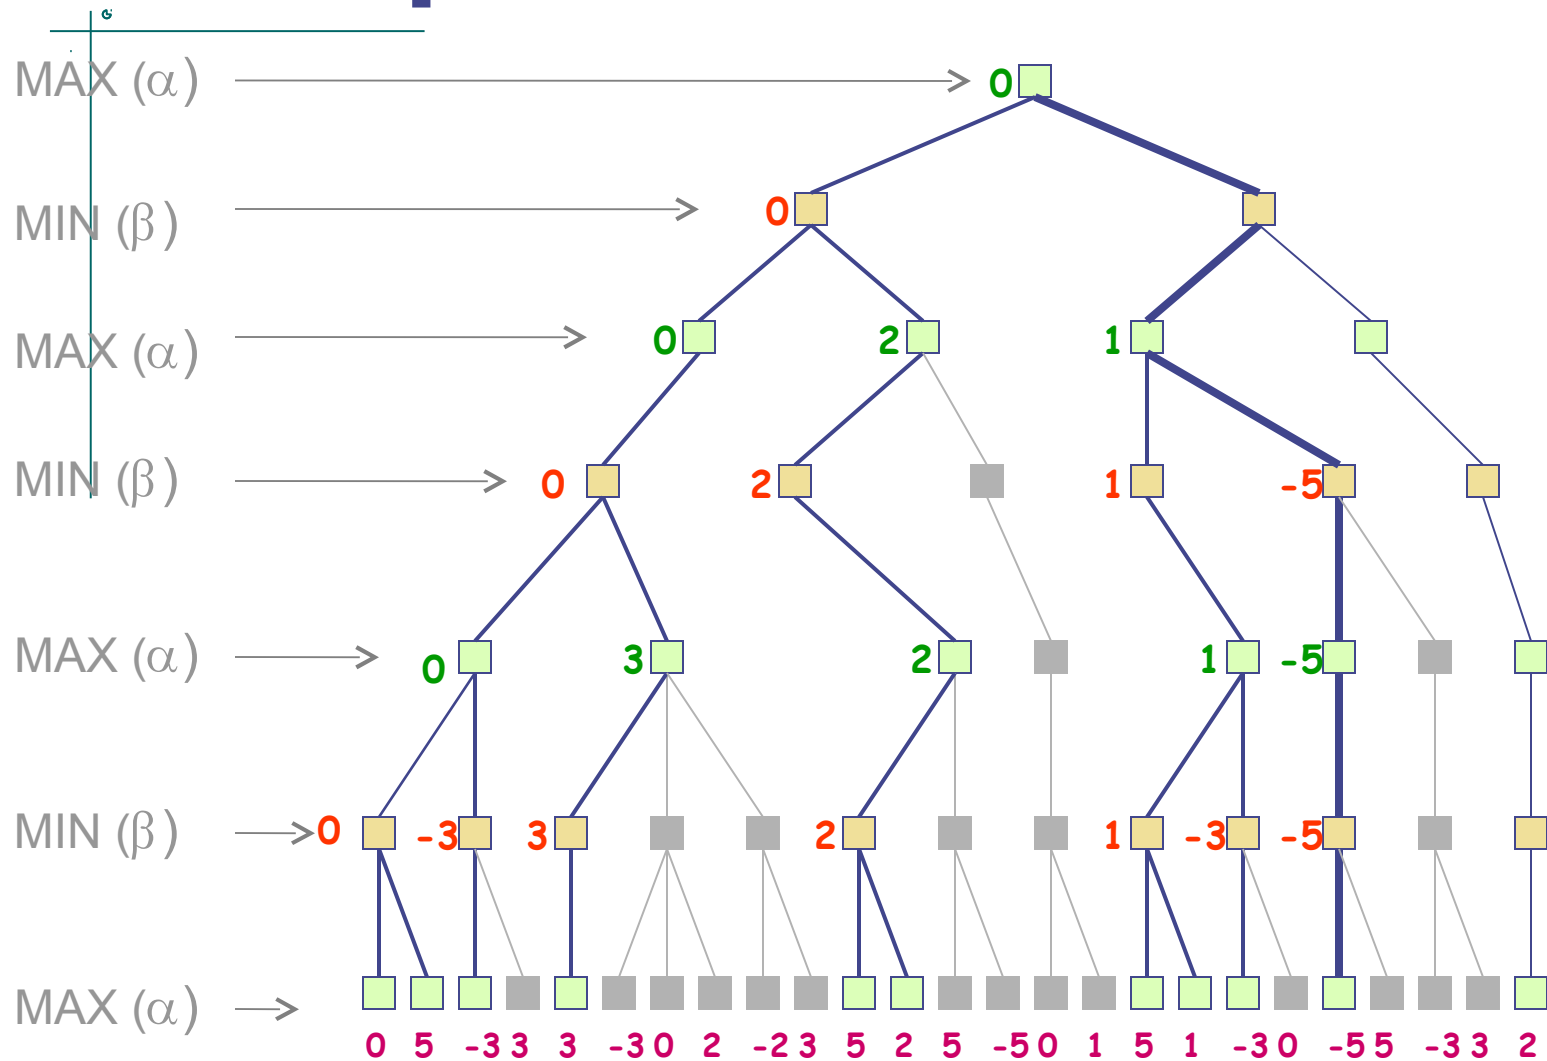
# Example



# Example

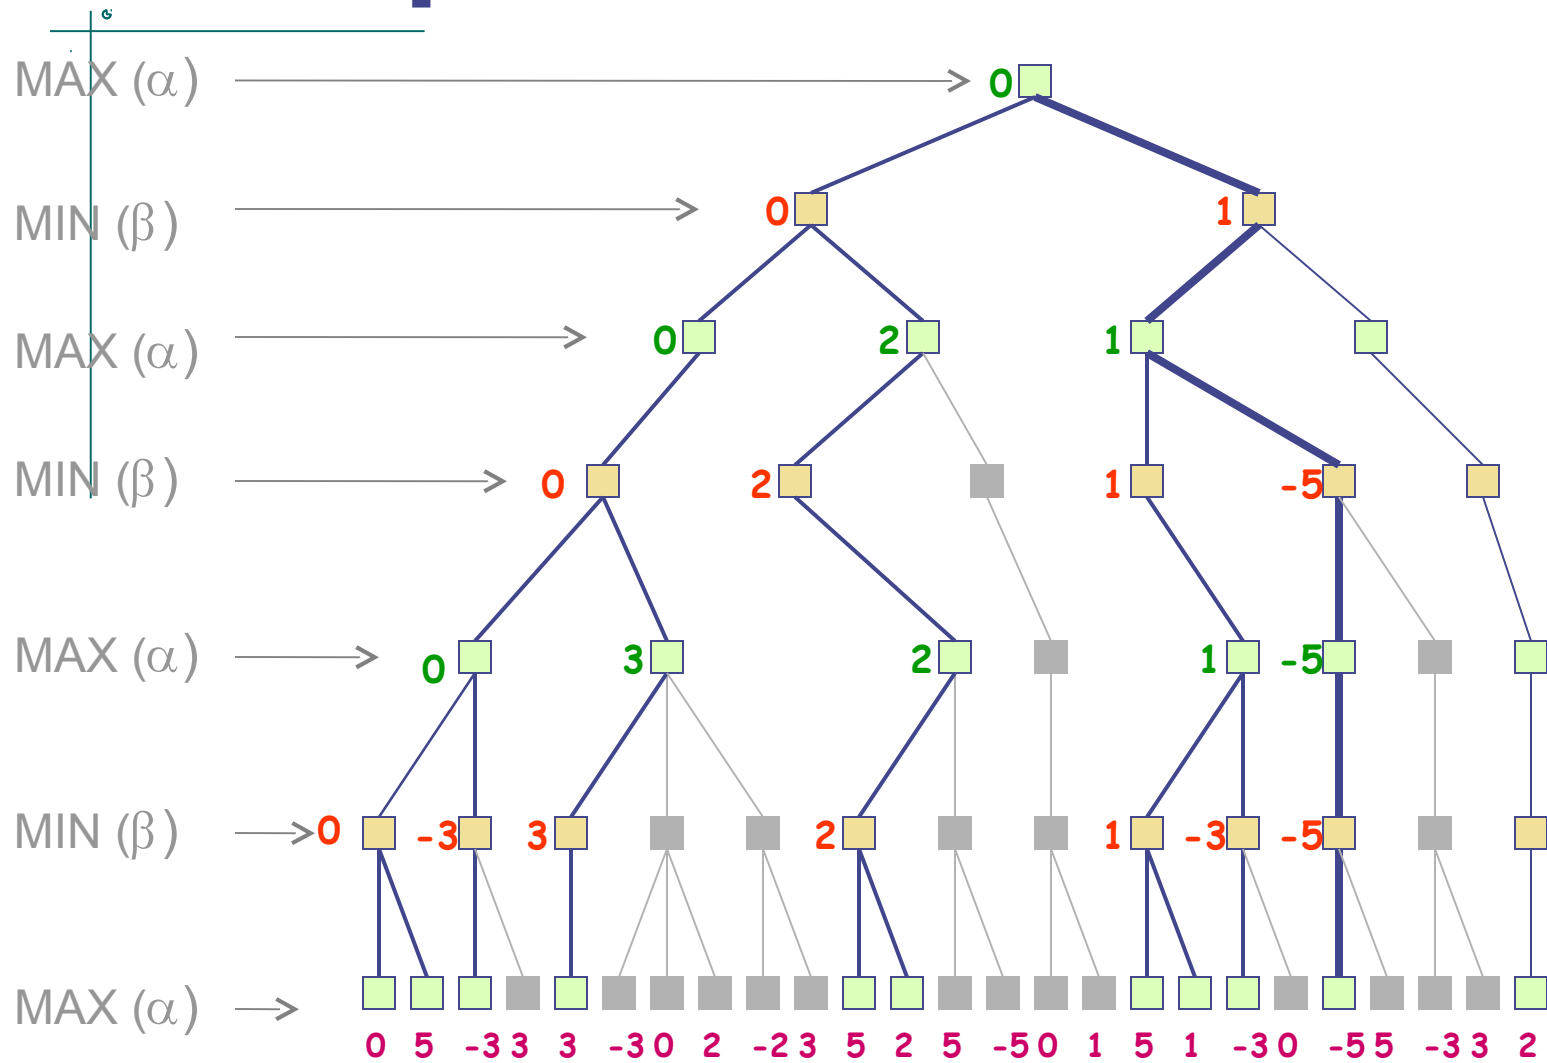


# Example

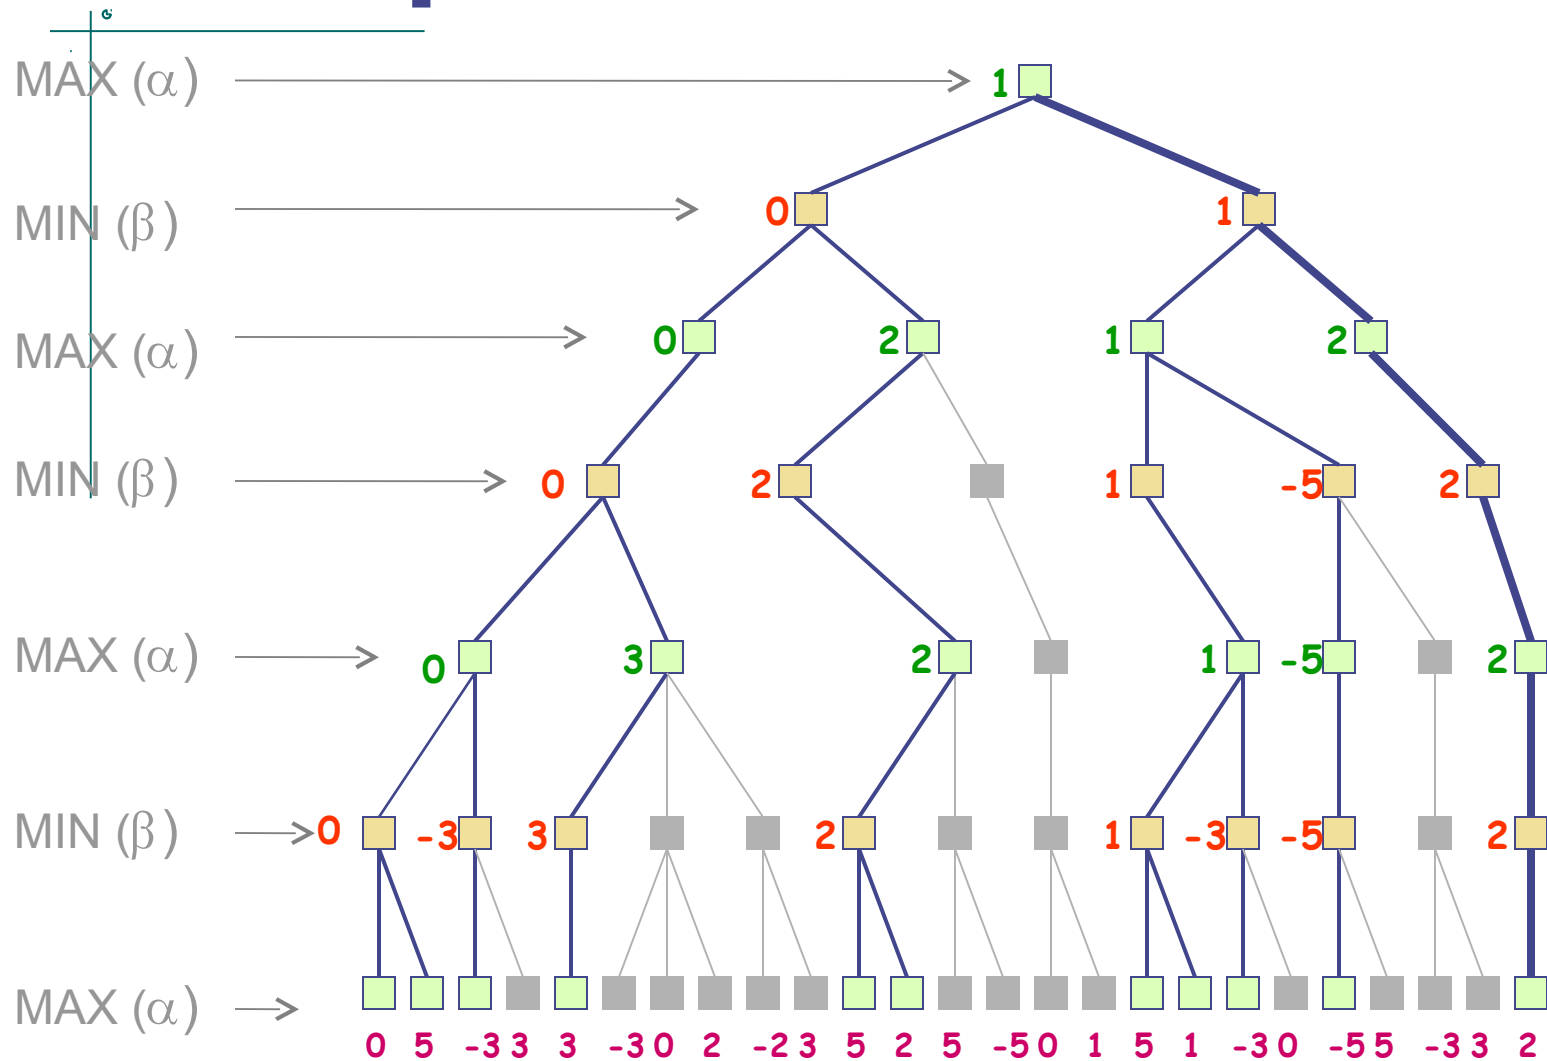




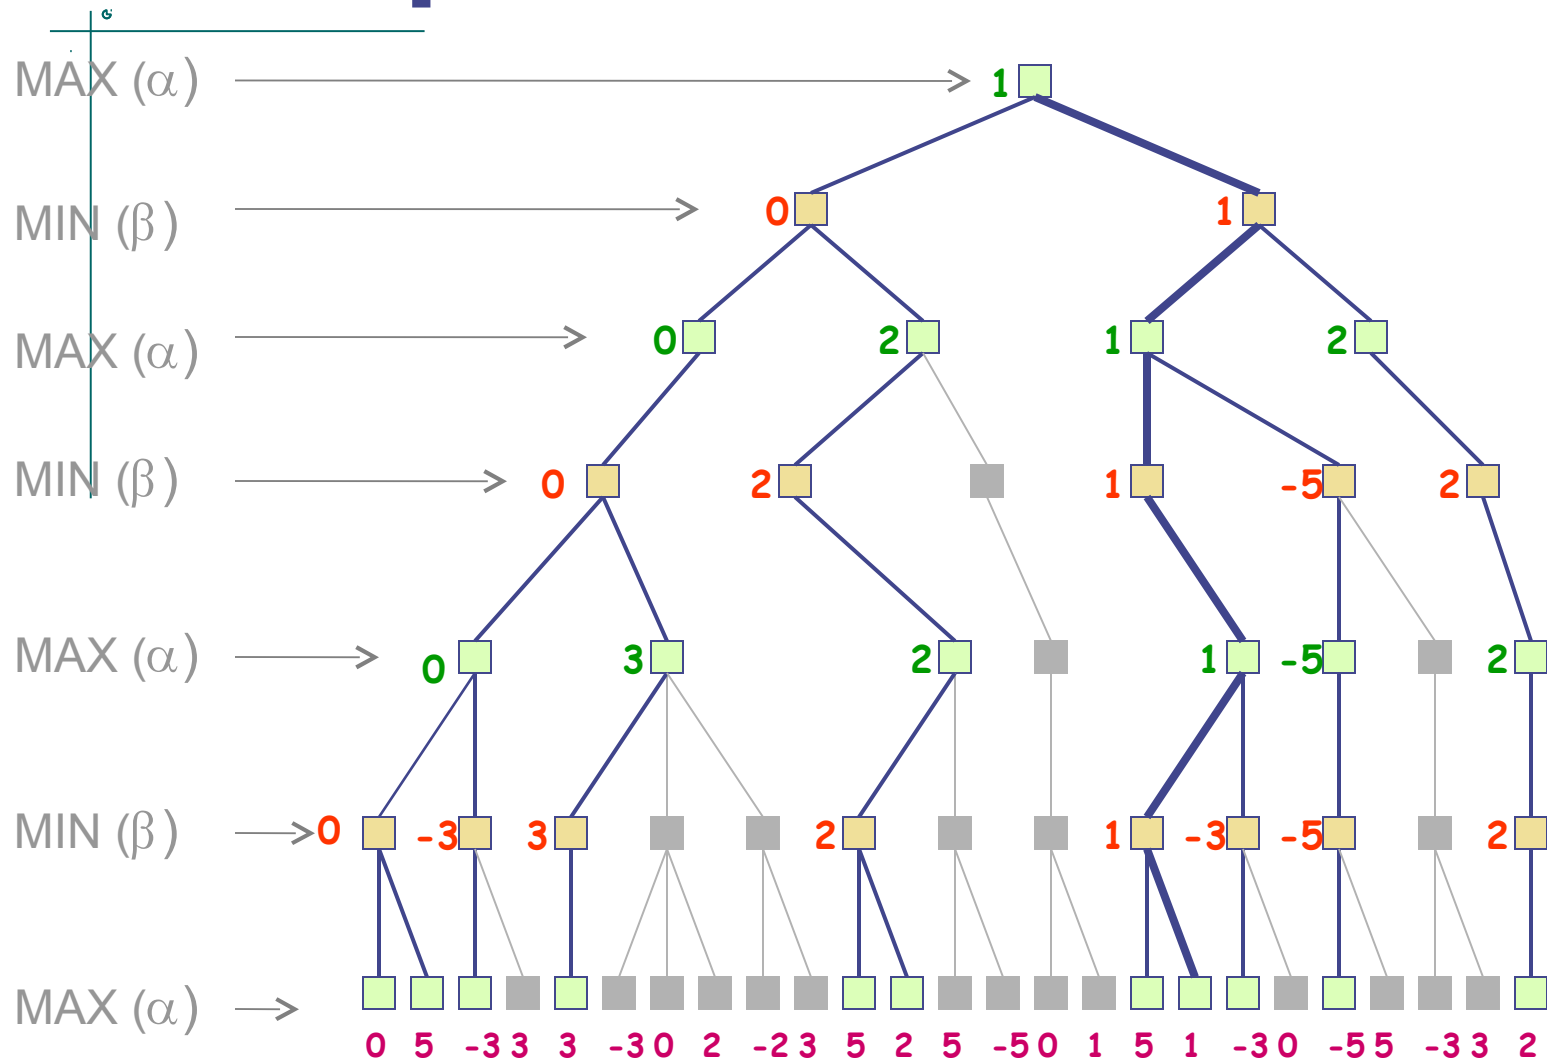
# Example



# Example





# Example



# How much do we gain?

Consider these two cases:



# How much do we gain?

- Assume a game tree of uniform branching factor  $b$
- Minimax examines  $O(b^h)$  nodes, so does alpha-beta in the worst-case

# How much do we gain?

- The **gain** for alpha-beta is maximum when:
  - The **MIN children** of a MAX node are ordered in **decreasing** backed up values
  - The **MAX children** of a MIN node are ordered in **increasing** backed up values
- Then alpha-beta examines  $O(b^{h/2})$  nodes [Knuth and Moore, 1975]

# How much do we gain?

- **But** this requires an oracle (if we knew how to order nodes perfectly, we would not need to search the game tree)
- If nodes are **ordered at random**, then the average number of nodes examined by alpha-beta is  $O(b^{3h/4})$ 
  - Good move ordering is essential for efficient alpha-beta pruning!

# Heuristic Ordering of Nodes

- Use iterative deepening
- Order the nodes below the root according to the values backed-up at the previous iteration



# Other Improvements

- Other heuristics to increase cut-offs:
  - Killer move heuristics / refutation table
  - Null-move heuristics
  - Null-window search / Negascout
- Use transposition tables to deal with repeated states
- To mitigate the horizon effect:  
Quiescence search/ Singular extension
- Forward pruning
- End-game databases
- Opening books


# Computers For The Win!

# Chinook (1994)



First computer to become official world champion of Checkers!

# Chess: Kasparov vs. Deep Blue




## Kasparov

## Deep Blue

5'10"	Height	6' 5"
176 lbs	Weight	2,400 lbs
34 years	Age	4 years
50 billion neurons	Computers	32 RISC processors + 256 VLSI chess engines
2 pos/sec	Speed	200,000,000 pos/sec
Extensive	Knowledge	Primitive
Electrical/chemical	Power Source	Electrical
Enormous	Ego	None

1997: Deep Blue wins by 3 wins, 1 loss, and 2 draws

# Othello: Murakami vs. Logistello



Takeshi Murakami  
World Othello Champion

1997: The Logistello software crushed Murakami  
by 6 games to 0

# Go: Goemate vs. ??



Name: Chen Zhixing

Profession: Retired

Computer skills:


self-taught programmer

Author of Goemate (arguably the best  
Go program available today)



Gave Goemate a 9 stone  
handicap and still easily  
beat the program,  
thereby winning \$15,000

# Go: Goemate vs. ??



Name: Chen Zhixing

Profession: Retired

Computer skills:

Go has too high a branching factor for existing search techniques

Current and future software must rely on huge databases and pattern-recognition techniques



gave Goemate a 5 stone handicap and still easily beat the program, thereby winning \$15,000

# Secrets

Many game programs are based on alpha-beta + iterative deepening + extended/singular search + transposition tables + huge databases + ...

For instance, Chinook searched all checkers configurations with 8 pieces or less and created an endgame database of 444 billion board configurations



# Secrets

The methods are general, but their implementation is dramatically improved by many specifically tuned-up enhancements (e.g., the evaluation functions)

# Other Types of Games

- Multi-player games, with alliances or not
- Games with randomness in successor function (e.g., rolling a dice)
  - **Expectiminimax** algorithm (adds chance nodes to the Minimax tree)
- Games with partially observable states (e.g., card games)
  - Search of belief state spaces