

# Reinforcement Learning

Author: Stephan Schiffel

Adapted by Stefán Ólafsson


# Outline

- 1 What is RL
- 2 Multi-armed Bandits
- 3 Markov Decision Processes
- 4 Monte-Carlo Method

# Reinforcement Learning

- Idea: an agent is learning from the interactions with an environment
- What to learn: utility of states, which actions to do
- Goal: maximizing (long-term) pay-off
- Observations: states and reward/cost of actions

# Agent-Environment Interface



Agent ...

- interacts with environment at discrete time steps  $t = 0, 1, 2, \dots$
- observes state  $S_t$  and responds with action  $A_t$
- observed resulting reward  $R_{t+1}$  and state  $S_{t+1}$

# Challenges of RL

- Evaluative feedback (reward)
- Delayed consequences / feedback
- Need for trial and error / exploration and exploitation
- Non-stationary processes

# Some Successes of RL

- World's best Backgammon player (Tesauro 1995)
- Acrobatic helicopter autopilots (Ng, Abbeel, Coates et.al. 2006)
- Widely used in placement and selection of ads on the web
- Strategic decision making in Jeopardy (IBM's Watson 2011)
- Human-level performance on Atari games from pixel-level input (Google Deepmind 2015)
- Super human-level play in Go and Chess (AlphaGo, AlphaZero 2017)
- Self-play to learn *Diplomacy* (Cicero 2022)
- RL with human feedback (ChatGPT 2022)

# Outline

- 1 What is RL
- 2 Multi-armed Bandits**
- 3 Markov Decision Processes
- 4 Monte-Carlo Method

# Multi-armed Bandits

- One-armed Bandit = slot machine



- Suppose you are in a casino and have a multiple slot machines  
⇒ Multi-armed Bandit
- Goal: maximize your profits



# Examples

- Action 1 – Reward is always 8  
value of action 1 is  $q_*(1) =$
- Action 2 – 88% chance of 0, 12% chance of 100  
 $q_*(2) =$
- Action 3 – reward is uniformly random between -10 and 35  
 $q_*(3) =$
- Action 4 – a third 0, a third 21, and a third from  $\{8, 9, \dots, 18\}$   
 $q_*(4) =$

# Examples

- Action 1 – Reward is always 8  
value of action 1 is  $q_*(1) = 8$
- Action 2 – 88% chance of 0, 12% chance of 100  
 $q_*(2) =$
- Action 3 – reward is uniformly random between -10 and 35  
 $q_*(3) =$
- Action 4 – a third 0, a third 21, and a third from  $\{8, 9, \dots, 18\}$   
 $q_*(4) =$

# Examples

- Action 1 – Reward is always 8  
value of action 1 is  $q_*(1) = 8$
- Action 2 – 88% chance of 0, 12% chance of 100  
 $q_*(2) = 0.88 \times 0 + 0.12 \times 100 = 12$
- Action 3 – reward is uniformly random between -10 and 35  
 $q_*(3) =$
- Action 4 – a third 0, a third 21, and a third from  $\{8, 9, \dots, 18\}$   
 $q_*(4) =$

# Examples

- Action 1 – Reward is always 8  
value of action 1 is  $q_*(1) = 8$
- Action 2 – 88% chance of 0, 12% chance of 100  
 $q_*(2) = 0.88 \times 0 + 0.12 \times 100 = 12$
- Action 3 – reward is uniformly random between -10 and 35  
 $q_*(3) = \sum_{i=-10}^{35} \frac{1}{46} \times i = \frac{-10+35}{2} = 12.5$
- Action 4 – a third 0, a third 21, and a third from  $\{8, 9, \dots, 18\}$   
 $q_*(4) =$

# Examples

- Action 1 – Reward is always 8  
value of action 1 is  $q_*(1) = 8$
- Action 2 – 88% chance of 0, 12% chance of 100  
 $q_*(2) = 0.88 \times 0 + 0.12 \times 100 = 12$
- Action 3 – reward is uniformly random between -10 and 35  
 $q_*(3) = \sum_{i=-10}^{35} \frac{1}{46} \times i = \frac{-10+35}{2} = 12.5$
- Action 4 – a third 0, a third 21, and a third from  $\{8, 9, \dots, 18\}$   
 $q_*(4) = \frac{1}{3} \times 21 + \frac{1}{3} \times \frac{8+18}{2} = 20$

# The k-armed Bandit Problem

- Infinite Sequence of time steps  $t = 1, 2, 3, \dots$
- Each step you choose an action  $A_t$  from  $k$  possibilities and receive a reward  $R_t$
- Reward is probabilistic and comes from some unknown (but fixed) distribution

$$q_*(a) = \mathbb{E}[R_t | A_t = a], \forall a \in 1, \dots, k$$

- True values  $q_*(a)$  are unknown
- Still, you must maximize the total reward
- You must both try actions to learn their values (explore) and prefer those that appear best (exploit)

# Exploration vs. Exploitation

- Suppose you learn estimates

$$Q_t(a) \approx q_*(a), \forall a$$

- For example, estimate action values as sample-averages:

$$Q_t(a) = \frac{\text{sum of rewards when } a \text{ was taken}}{N_t(a)}$$

where  $N_t(a)$  is number of times  $a$  was taken before time  $t$ .

- Sample-average estimates converge to true values:

$$\lim_{N_t(a) \rightarrow \infty} Q_t(a) = q_*(a)$$

- **Greedy action** at time  $t$ :  $A_t^* = \operatorname{argmax}_a Q_t(a)$
- **Exploitation** = picking the greedy action:  $A_t = A_t^*$
- **Exploration** = picking a different action:  $A_t \neq A_t^*$

# Selection Strategy: $\epsilon$ -Greedy

- Mostly select the greedy action (with probability  $1 - \epsilon$ ), but
- with small probability  $\epsilon$  select a random action
- Perhaps the simplest way to balance exploration and exploitation



# Selection Strategy: Upper Confidence Bound (UCB)

- A clever way of reducing exploration over time
- Estimate an upper bound on the true action value
- Select the action with the highest upper bound (optimistic estimated value)

$$A_t = \underset{a}{\operatorname{argmax}} \left[ Q_t(a) + c \sqrt{\frac{\log t}{N_t(a)}} \right]$$

# Update Rule for $Q$

- Goal: update the action value estimate one step at a time (and forget the details about the previous time steps)
- Keeping a running average:

$$Q_{t+1}(a) = \begin{cases} Q_t(a) + \frac{1}{N_t(a)} [R_t - Q_t] & \text{if } A_t = a \\ Q_t(a) & \text{otherwise} \end{cases}$$

$$N_{t+1}(a) = \begin{cases} N_t(a) + 1 & \text{if } A_t = a \\ N_t(a) & \text{otherwise} \end{cases}$$

- General update rule:

$$\text{NewEstimate} = \text{OldEstimate} + \text{StepSize} [\text{Target} - \text{OldEstimate}]$$

# Tracking a Non-Stationary Problem

- Suppose the true action values (and the reward distribution) change slowly over time
- Then we say the problem is non-stationary
- In this case sample-averages are not a good idea (Why?)
- Better is an exponential, recency-weighted average:

$$Q_{n+1} = Q_n + \alpha * [R_n - Q_n]$$

where **step size**  $\alpha$  is a constant  $0 \leq \alpha \leq 1$

- There is bias due to  $Q_1$  that becomes smaller over time.

# Outline

- 1 What is RL
- 2 Multi-armed Bandits
- 3 Markov Decision Processes**
- 4 Monte-Carlo Method

# Markov Decision Process

- ... is a (typically non-deterministic) state transition system
- Defined by transition probabilities:

$$p(s', r | s, a) = P(S_{t+1} = s', R_{t+1} = r | S_t = s, A_t = a)$$

- Transition probabilities between states depend on agent's action
- State transitions yield rewards
- Markov property: The future only depends on current state and the agent's future actions, not on past states or actions

# Goal: Learning a Policy

- Policy  $\pi$  is a mapping from states to action probabilities

$$\pi(a|s) = \text{probability that } A_t = a \text{ if } S_t = s$$

- Special case: deterministic policy

$$\pi(s) = \text{the action taken with prob. 1 in state } s$$

- Reinforcement Learning methods specify how  $\pi$  changes over time as a result of the agent's experience
- Typically,  $\pi(a|s)$  is defined in relation to  $Q(s, a)$
- Goal: Learning a policy that increases reward in the long run

# Return = cumulative reward over time

- $G_t \dots$  return from time  $t$  onwards

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

where  $\gamma$ ,  $0 \leq \gamma \leq 1$  is the **discount rate**.

- shortsighted  $0 \leftarrow \gamma \rightarrow 1$  farsighted
- If  $\gamma = 1$ , return is the cumulative future reward
- If  $\gamma = 0$ , return is only the immediate reward
- Typically,  $\gamma \in \{0.9, 0.99, \dots\}$

## 4 Value Functions

- $v_{\pi}(s)$  ... expected return in state  $s$  if all future actions are picked with policy  $\pi$
- $q_{\pi}(s, a)$  ... expected return in state  $s$  if the first action is  $a$  and afterwards actions are picked with policy  $\pi$
- $v_{*}(s) = \max_{\pi} v_{\pi}(s)$  (optimal value of a state)
- $q_{*}(s, a) = \max_{\pi} q_{\pi}(s, a)$  (optimal value of a state-action pair)
- All of these are theoretic values, in practice we only have their estimates  $V_t(s)$  and  $Q_t(s, a)$



## 4 Value Functions

- $v_\pi(s)$  ... expected return in state  $s$  if all future actions are picked with policy  $\pi$
- $q_\pi(s, a)$  ... expected return in state  $s$  if the first action is  $a$  and afterwards actions are picked with policy  $\pi$
- $v_*(s) = \max_\pi v_\pi(s)$  (optimal value of a state)
- $q_*(s, a) = \max_\pi q_\pi(s, a)$  (optimal value of a state-action pair)
- All of these are theoretic values, in practice we only have their estimates  $V_t(s)$  and  $Q_t(s, a)$
- Optimal policy:  $\pi_*$  is optimal if it only picks optimal actions, i.e.

$$\pi_*(a|s) > 0 \text{ only when } q_*(s, a) = \max_b q_*(s, b)$$

- In other words: A policy is optimal iff it is greedy wrt.  $q_*$ .

# Bellman Equation for State Values

$$v_{\pi}(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_{\pi}(s')]$$

- In case of finite MDPs (finitely many states and actions), this is a set of (linear) equations.
- The unique solution is  $v_{\pi}$ .

# Bellman Equation for State Values

$$v_{\pi}(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_{\pi}(s')]$$

- In case of finite MDPs (finitely many states and actions), this is a set of (linear) equations.
- The unique solution is  $v_{\pi}$ .
- Similarly,

$$v_*(s) = \max_a \sum_{s', r} p(s', r|s, a) [r + \gamma v_*(s')]$$

- However, these equations are non-linear.

# Bellman Equation for Action Values

- Note that,

$$v_{\pi}(s) = \sum_a \pi(a|s) q(s, a)$$

and

$$q_{\pi}(s, a) = \sum_{s', r} p(s', r|s, a) [r + \gamma v_{\pi}(s')]$$

- Therefore,

$$q_{\pi}(s, a) = \sum_{s', r} p(s', r|s, a) \left[ r + \gamma \sum_{a'} \pi(a'|s') q_{\pi}(s', a') \right]$$

and

$$q_*(s, a) = \sum_{s', r} p(s', r|s, a) \left[ r + \gamma \max_{a'} q_*(s', a') \right]$$

# Outline

- 1 What is RL
- 2 Multi-armed Bandits
- 3 Markov Decision Processes
- 4 Monte-Carlo Method**

# Monte-Carlo Methods

- Idea: observe a whole episode (sequence of state transitions and rewards) and learn  $Q_t(s, a)$  from the observed return  $G$
- Define policy (somewhat) greedy wrt.  $Q_t(s, a)$  and generate more episodes to learn from  
 $\Rightarrow$  simultaneously learns  $Q(s, a)$  and  $\pi(a|s)$
- Converges to (something close to)  $\pi_*$  under reasonable constraints
- No model necessary (transition probabilities can be unknown)

# Policy Evaluation (Estimating $v_\pi$ )

- Repeatedly, generate episodes starting in state  $s$  and observe the return (cumulative discounted reward)  $G$
- For every observation  $(s, G)$  update  $V(s)$  with

$$V(s) \leftarrow V(s) + \alpha [G - V(s)]$$

- $V(s)$  will eventually converge to  $v_\pi(s)$ , as long as all states  $s$  are visited infinitely often (and  $\alpha \rightarrow 0$  with  $t \rightarrow \infty$ ).

# Policy Evaluation (Estimating $v_\pi$ )

- Repeatedly, generate episodes starting in state  $s$  and observe the return (cumulative discounted reward)  $G$
- For every observation  $(s, G)$  update  $V(s)$  with

$$V(s) \leftarrow V(s) + \alpha [G - V(s)]$$

- $V(s)$  will eventually converge to  $v_\pi(s)$ , as long as all states  $s$  are visited infinitely often (and  $\alpha \rightarrow 0$  with  $t \rightarrow \infty$ ).
- Compare with the Bellman-Equation! Why does it make sense?

$$v_\pi(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_\pi(s')]$$



# Learning a Policy with Monte-Carlo

## (On-Policy Monte-Carlo Control)

- Initialize:  $Q(s, a)$  arbitrarily
- Let  $\pi(a|s)$  be  $\epsilon$ -greedy wrt.  $Q(s, a)$
- Repeat:
  - ▶ Generate an episode using  $\pi$
  - ▶ For each pair  $s, a$  in the episode
    - ★  $G$  is the return following  $s, a$
    - ★  $Q(s, a) \leftarrow Q(s, a) + \alpha [G - Q(s, a)]$
  - ▶ Let  $\pi$  be epsilon-greedy wrt. the updated  $Q(s, a)$
- $\Rightarrow \pi$  converges to best  $\epsilon$ -greedy policy (if  $\alpha = \frac{1}{N(s,a)}$ ).

# Monte-Carlo Tree Search

- Similar to on-policy Monte-Carlo control
- Only keeps track of  $Q(s, a)$  for selected states, but eventually learns  $Q(s, a)$  for all states
- Learns separate  $Q(s, a)$  and policies for all agents simultaneously
- Uses UCB instead of  $\epsilon$ -greedy, that is, it reduces exploration over time  $\Rightarrow$  ensures convergence to optimal policy (in case of perfect information games)

# Variation: Temporal Difference Learning

- Only observe single state transitions instead of whole episodes
- For every observation  $(s, s', r)$  update  $V(s)$  with

$$V(s) \leftarrow V(s) + \alpha [r + \gamma V(s') - V(s)]$$

- Works for continuous environments (infinitely long episodes)

# Variation: Function Approximation

- So far,  $Q(s, a)$  (or  $V(s)$ ) are represented as a lookup-table with an entry for each  $s, a$
- For interesting problems:
  - ▶ Too many states
  - ▶ Too slow to learn the value of each state individually
- Solution for large MDPs:
  - ▶ Estimate value function with function approximation

$$\hat{v}(s, \vec{\theta}) \approx V(s) \approx v_{\pi}(s) \quad (1)$$

$$\hat{q}(s, a, \vec{\theta}) \approx Q(s, a) \approx q_{\pi}(s, a) \quad (2)$$

- ▶ Generalize from seen states to unseen states
- ▶ Update  $\vec{\theta}$  (instead of  $V$  or  $Q$ ) using MC or TD learning

# Summary

- Agent-environment interface
- Multi-arm bandits
- Markov decision processes
- Bellman equation
- Monte-Carlo is a sampling-based solution of the Bellman equation interleaved with updates of the policy