



# Neural Networks

## Neural Language Models

T-622-ARTI

Stefán Ólafsson  
Spring 2023







# Neural Network

- A network of small computing units (nodes)
- Each unit takes a vector of input values
- Produces a single output value

Single unit



Network





# Feed-forward network and deep learning

- **Feed-forward:**
  - Computation proceeds iteratively from one layer of units to the next
- **Deep learning:**
  - Networks that are deep (many layers)



# Neural Network (NN) and logistic regression


- **NN share much of the same mathematics as logistic regression**
  - But a more powerful classifier than logistic regression
  - Because of these hidden layers
- **Logistic regression uses feature templates based on domain knowledge**
- **More common in NN to avoid the use of hand-derived features**
  - Take raw words as inputs
  - Learns to induce features as part of the learning process



# Units

- A neural unit takes a **weighted sum** of its input and adds a **bias term**
- Input:  $x_1, x_2 \dots x_n$
- Weights:  $w_1, w_2 \dots w_n$
- $z = \sum_{i=1}^n w_i x_i + b$
- $z = w * x + b$

Single unit





# Activation value


- Neural units apply a non-linear function  $f$  to  $z$  to produce the output, the activation value,  $a$ .

$$y = a = f(z)$$


- Popular non-linear functions: sigmoid, tanh, ReLU




**(a)**  $\sigma(x) = 1/(1 + e^{-x})$



**(b)**  $\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$




**(c)**  $f(x) = \max(0, x)$





# Example



**Figure 7.2** A neural unit, taking 3 inputs  $x_1$ ,  $x_2$ , and  $x_3$  (and a bias  $b$  that we represent as a weight for an input clamped at  $+1$ ) and producing an output  $y$ . We include some convenient intermediate variables: the output of the summation,  $z$ , and the output of the sigmoid,  $a$ . In this case the output of the unit  $y$  is the same as  $a$ , but in deeper networks we'll reserve  $y$  to mean the final output of the entire network, leaving  $a$  as the activation of an individual node.





# Why do we need non-linear activation functions?

- to introduce **non-linearity** into the network
- this allows you to model output that varies non-linearly with its input variables
- if we only allow linear activation functions in a neural network, the output will just be a linear transformation of the input
- non-linear means that the output cannot be reproduced from a linear combination of the input



# Why do we need multi-layer networks?

- **Because a single neural unit cannot compute some very simple functions of its input.**



# Logical AND and OR



**Figure 7.4** The weights  $w$  and bias  $b$  for perceptrons for computing logical functions. The inputs are shown as  $x_1$  and  $x_2$  and the bias as a special node with value  $+1$  which is multiplied with the bias weight  $b$ . (a) logical AND, showing weights  $w_1 = 1$  and  $w_2 = 1$  and bias weight  $b = -1$ . (b) logical OR, showing weights  $w_1 = 1$  and  $w_2 = 1$  and bias weight  $b = 0$ . These weights/biases are just one from an infinite number of possible sets of weights and biases that would implement the functions.

$$y = \begin{cases} 0, & \text{if } w * x + b \leq 0 \\ 1, & \text{if } w * x + b > 0 \end{cases}$$

**Perceptron:** binary classifier;  
purely linear




# Logical XOR


- **Not possible to build a perceptron to compute logical XOR!**
- **A perceptron is a linear classifier**
- **XOR is not a linearly separable function**




# Linearly separable



a)  $x_1$  AND  $x_2$



b)  $x_1$  OR  $x_2$




c)  $x_1$  XOR  $x_2$

**Figure 7.5** The functions AND, OR, and XOR, represented with input  $x_0$  on the x-axis and input  $x_1$  on the y axis, Filled circles represent perceptron outputs of 1, and white circles perceptron outputs of 0. There is no way to draw a line that correctly separates the two categories for XOR. Figure styled after [Russell and Norvig \(2002\)](#).



# XOR solution




**Figure 7.6** XOR solution after [Goodfellow et al. \(2016\)](#). There are three ReLU units, in two layers; we've called them  $h_1$ ,  $h_2$  ( $h$  for “hidden layer”) and  $y_1$ . As before, the numbers on the arrows represent the weights  $w$  for each unit, and we represent the bias  $b$  as a weight on a unit clamped to +1, with the bias weights/units in gray.

**You should work out what happens for the inputs!**



# New representation



**Figure 7.7** The hidden layer forming a new representation of the input. Here is the representation of the hidden layer,  $h$ , compared to the original input representation  $x$ . Notice that the input point  $[0\ 1]$  has been collapsed with the input point  $[1\ 0]$ , making it possible to linearly separate the positive and negative cases of XOR. After [Goodfellow et al. \(2016\)](#).



# Hidden layer

- **Key advantage of neural networks:**
  - Can automatically learn useful representations of the input






# Feed-forward Neural Networks

- **Simple feed-forward networks have three kinds of nodes:**
  - Input units
  - Hidden units
  - Output units




# Feed-forward Neural Networks



**Figure 7.8** A simple 2-layer feed-forward network, with one hidden layer, one output layer, and one input layer (the input layer is usually not counted when enumerating layers).



# Deep Feed-forward Neural Network





# Feed-forward Neural Networks

- Weight matrix **W** for a hidden layer
- $W_{ij}$ : the weight of the connection from the  $i$ -th input  $x_i$  to the  $j$ -th hidden unit  $h_j$
- Allows for effective matrix operations
- $h = \sigma(W * x + b)$
- **W** is a matrix, **x**, **b** and **h** are vectors



# Output layer

- $\mathbf{h}$  forms a representation of the input
- The output layer takes this new representation  $\mathbf{h}$  and computes a final output.
- The output could be a real value
- In many cases, however, the network makes a classification decision.



# Output layer

- If binary classification, then single output node
- $y$  is then the probability of the positive class
- If multinomial classification (e.g. PoS tagging), then one output node for each possible PoS
- The output layer then gives a probability distribution over output nodes.




# Output layer

- The output layer has a weight matrix, **U**.
- $z = U * h$
- $z$  is a vector of real-valued numbers
- For multiclass classification we need a vector of probabilities
- Can convert it to a probability distribution by using the **softmax** function ( $d$  is the dimension of  $z$ ):
- $$\text{softmax}(z_i) = \frac{e^{z_i}}{\sum_{j=1}^d e^{z_j}} \quad 1 \leq i \leq d$$



# Feed-forward Neural Networks



**Figure 7.8** A simple 2-layer feed-forward network, with one hidden layer, one output layer, and one input layer (the input layer is usually not counted when enumerating layers).





# Softmax example

- **Given a vector**  $z=[0.6, 1.1, -1.5, 1.2, 3.2, -1.1]$
- **softmax( $z$ )** is  $[0.055, 0.090, 0.0067, 0.10, 0.74, 0.010]$ .



# Neural network classifier with one hidden layer

- Builds a vector  $\mathbf{h}$ , a hidden layer representation of the input
- Then runs standard logistic regression on the features that the network develops in  $\mathbf{h}$ .
- Deep neural network is like layer after layer of logistic regression classifiers:
  - Prior layers induce the feature representations themselves, as opposed to using hand-crafted features



# Neural network classifier with one hidden layer

- $h = \sigma(W * x + b)$
- $z = U * h$
- $y = \text{softmax}(z)$
- 2-layer network



# Training

- Learn weights  $\mathbf{W}^{[l]}$  and bias  $\mathbf{b}^{[l]}$  that make  $\hat{\mathbf{y}}$  for each training observation as close as possible to the true  $\mathbf{y}$ .
1. loss (cost) function
  2. Iteratively updating the weights => gradient descent




# Difference to logistic regression?

- In logistic regression, for each observation we directly compute the derivative of the loss function with respect to an individual **w** or **b**.
  - Cost function is convex
- For neural networks, with many layers, we need what is called **error back-propagation**
- See chapter 7.4 in the textbook



# More than one layer: non-convex





# Application: Neural Language Models (NLM)

- **Language modeling:**
  - Predicting upcoming words from prior word contexts
- **Advantages of NLMs compared to n-gram language models:**
  - Don't need smoothing
  - Can handle much longer histories
  - Can generalize over context of similar words
  - Usually give higher accuracy



# Neural Language Models (NLM)

- **Disadvantages of NLMs compared to n-gram language models:**
  - Strikingly slower to train





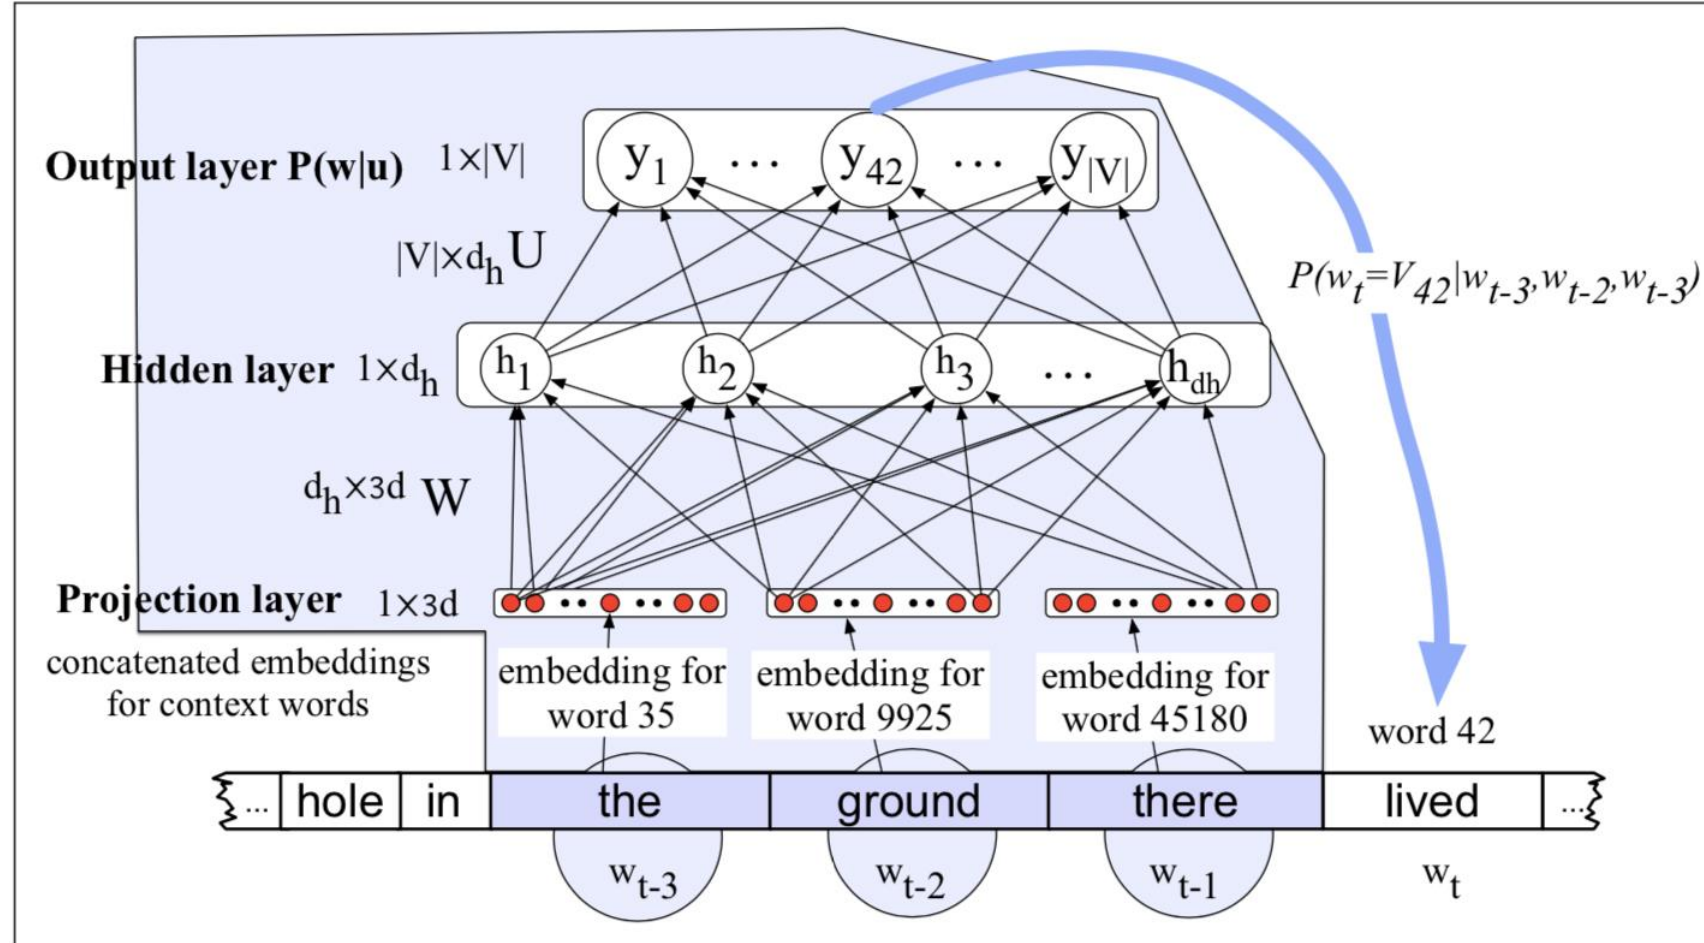
# Feed-forward neural language model

- Standard feedforward network
- Takes as input at time  $t$  a representation of some number of previous words ( $w_{t-1}, w_{t-2} \dots$ )
- Output a probability distribution over possible next words
- $P(w_t | w_1^{t-1}) \approx P(w_t | w_{t-N+1}^{t-1})$
- $N=4$ , 4-gram:  $P(w_t | w_{t-1}, w_{t-2}, w_{t-3})$



# Word Embeddings

- Prior context is represented by embeddings of the previous words.
- Allows neural language models to generalize to unseen data, much better than n-gram LMs




**Figure 7.12** A simplified view of a feedforward neural language model moving through a text. At each timestep  $t$  the network takes the 3 context words, converts each to a  $d$ -dimensional embeddings, and concatenates the  $3 \times Nd$  unit input layer  $x$  for the network. These units are multiplied by a weight matrix  $W$  and bias vector  $b$  and then an activation function to produce a hidden layer  $h$ , which is then multiplied by another weight matrix  $U$ . (For graphic simplicity we don't show  $b$  in this and future pictures). Finally, a softmax output layer predicts at each node  $i$  the probability that the next word  $w_t$  will be vocabulary word  $V_i$ . (This picture is simplified because it assumes we just look up in an embedding dictionary  $E$  the  $d$ -dimensional embedding vector for each word, precomputed by an algorithm like word2vec.)



# What if we do not have pre-trained embeddings?

- We then need to learn the embeddings during training of the network
- We represent each of the previous words as a one-hot vector
  - one dimension for each word in the vocabulary
  - $[0\ 0\ 0\ 0\ 1\ 0\ 0\ \dots\ 0\ 0\ 0\ 0]$



**Figure 7.13** Learning all the way back to embeddings. Notice that the embedding matrix  $E$  is shared among the 3 context words.